

Advanced Spark Operations

Rosa Filgueira

Working with Key-Value Pairs

- Spark's “distributed reduce” transformations act on RDDs of *key-value pairs*
- *Key-value* RDDs let you:
 - *Count up reviews for each product*
 - *Group together data with the same key*
 - *Group together two different RDDs*
- Python:

```
pair = (a, b)
pair[0] # => a → key
pair[1] # => b → value
```



= easy



= medium

Essential Core & Intermediate PairRDD Operations

General

- flatMapValues
- groupByKey
- reduceByKey
- reduceByKeyLocally
- foldByKey
- aggregateByKey
- sortByKey
- combineByKey

Math / Statistical

- sampleByKey

Set Theory / Relational

- cogroup (=groupWith)
- join
- subtractByKey
- fullOuterJoin
- leftOuterJoin
- rightOuterJoin

Data Structure

- partitionBy

-
- keys
 - values

- countByKey
- countByValue
- countByValueApprox
- countApproxDistinctByKey
- countApproxDistinctByKey
- countByKeyApprox
- sampleByKeyExact

TRANSFORMATIONS

ACTIONS

Some Key-Value Transformations

```
pets = sc.parallelize([("cat", 1), ("dog", 1), ("cat", 2)])

pets.reduceByKey(lambda x, y: x + y)
# => {(cat, 3), (dog, 1)}

pets.groupByKey()
# => {(cat, 3), (dog, 1)}

pets.sortByKey() # by default ascending parameter
# => {(cat, 1), (cat, 2), (dog, 1)}

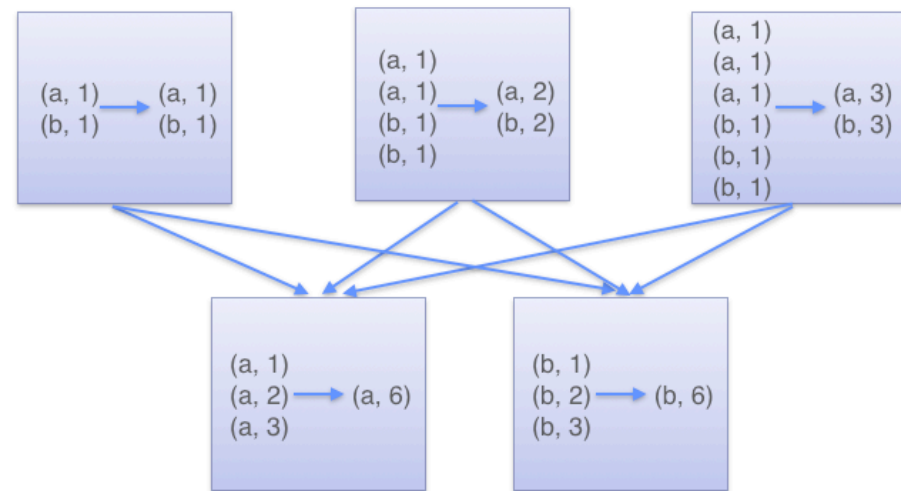
pets.map(lambda x: x[0], x[1] * 2)
# => {(cat, 2), (dog, 2), (cat, 4)}

pets.mapValues(lambda value: value * 2) # convenient to work only on
values
# => {(cat, 2), (dog, 2), (cat, 4)}

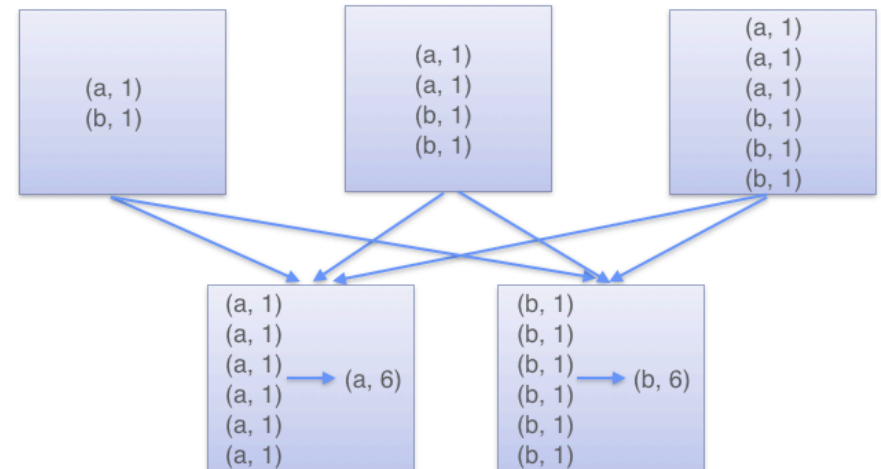
Pets.filter(lambda x: (x[1] > 1))
# => {(cat, 2)}
```

GroupByKey vs ReduceByKey

ReduceByKey



GroupByKey



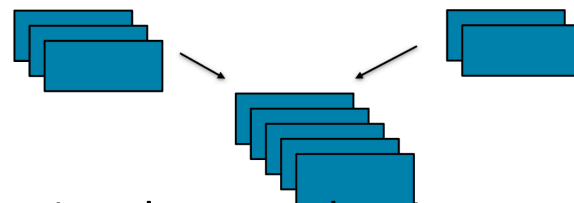
Both will give you the same answer:

- `reduceByKey` also automatically implements combiners on the map side
- `groupByKey` can cause a lot of data to be shuffled between workers
- `reduceByKey` is more efficient but has fixed memory space
- Prefer `reduceByKey`, `combineByKey`, `foldByKey` over `groupByKey`

UNION, JOIN (inner)

Union: Return a new RDD containing all items for two original RDDs. Duplicated are not culled

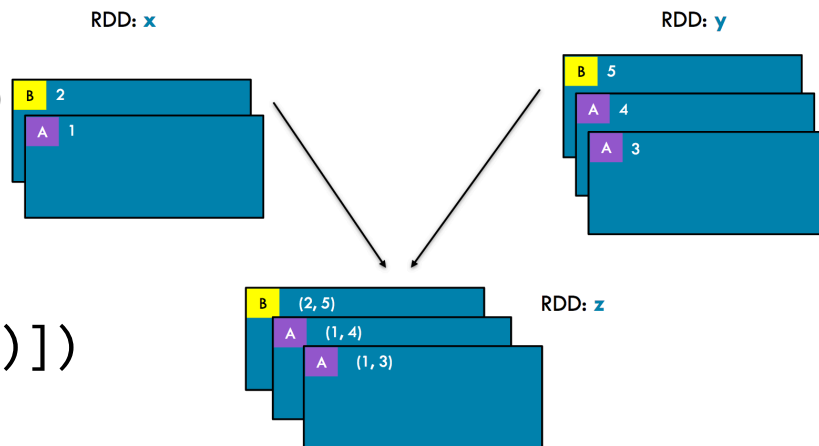
```
x= sc.parallelize([1,2,3])  
y= sc.parallelize([5,6])  
z= x.union(y) → [1,2,3,5,6]
```



Join: Return a new RDD containing all pairs of elements having the same key in the original RDD – inner join between 2 RDDs



```
x= sc.parallelize([("a",1), ("b,2")])  
y= sc.parallelize([("a",3), ("a,4"), ("b,5")])  
z= x.join(y) → [("a", (1,3), ("a", (1,4)),  
("b", (2,5))]
```



-- Another example

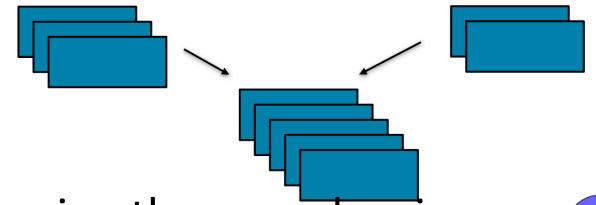
```
x= sc.parallelize([(1,2) (3,4), (3,6)])  
y= sc.parallelize([(3,9)])  
z = x.join(y) → ?
```



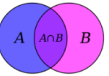
UNION, JOIN (inner)

Union: Return a new RDD containing all items for two original RDDs. Duplicated are not culled

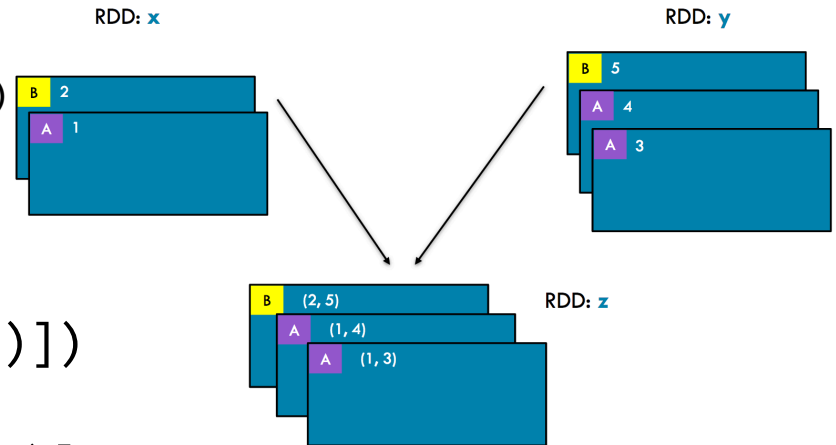
```
x= sc.parallelize([1,2,3])  
y= sc.parallelize([5,6])  
z= x.union(y) → [1,2,3,5,6]
```



Join: Return a new RDD containing all pairs of elements having the same key in the original RDD – inner join between 2 RDDs



```
x= sc.parallelize([("a",1), ("b,2")])  
y= sc.parallelize([("a",3), ("a,4"), ("b,5")])  
z= x.join(y) → [("a", (1,3), ("a", (1,4)),  
("b", (2,5))]
```



-- Another example

```
x= sc.parallelize([(1,2) (3,4), (3,6)])  
y= sc.parallelize([(3,9)])  
z = x.join(y) → [(3,(4,9)), (3, (6,9))]
```

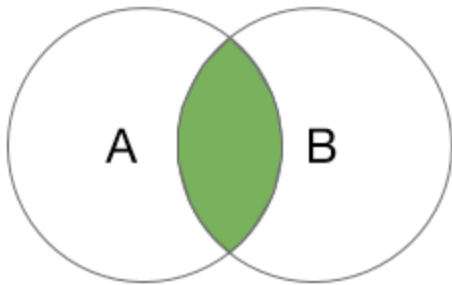


JOIN (outer)

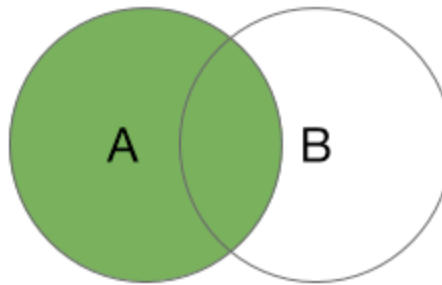
- Inner joins require a key to be present in both RDDs
- Outer joins do not require a key to be present in both RDDs:
 - leftOuterJoin
 - rightOuterJoin

```
x= sc.parallelize([(('A',1),('B',2), ('C',1))])
y= sc.parallelize([(('A',3), ('C',2), ('D',4))])
Z = x.leftOuter(x) → [(('A', (1,3)), ('C',(1,2)), ('B', (2,None)))]
Z = x.rightOuter(x) → [(('A', (1,3), ('C', (1,2), ('D', (4, None)))]
Z = x.join(x) → [(('A', (1,3), ('C', (1,2)))]
```

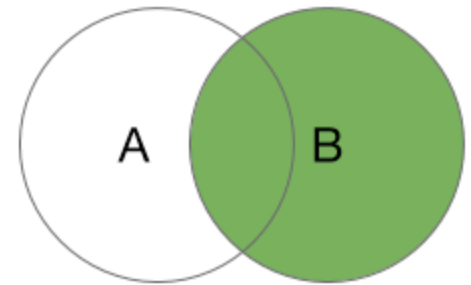

Beyond the traditional JOIN



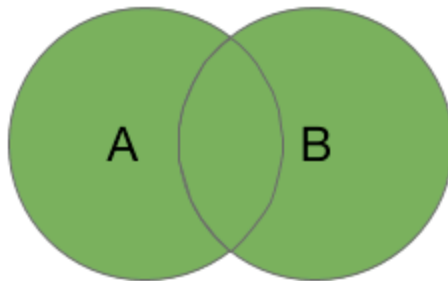
INNER JOIN



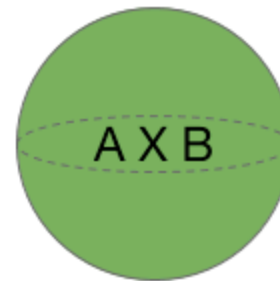
LEFT OUTER JOIN



RIGHT OUTER
JOIN



FULL OUTER
JOIN



CARTESIAN
(CROSS) JOIN

CoGroup

CoGroup: Given two RDDs sharing the same key type K, with the respective values types as V and W, the resulting RDD is of type [K, (iterable[V], Iterable[W])]

```
x= sc.parallelize([("A", (1,1)),("A", 2), ("B", (3,3)), ("C",4) ])
```

```
y= sc.parallelize([("D", (5,5)),("A", 6), ("B", 7), ("A",8) ])
```

```
z = x.cogroup(y)
```

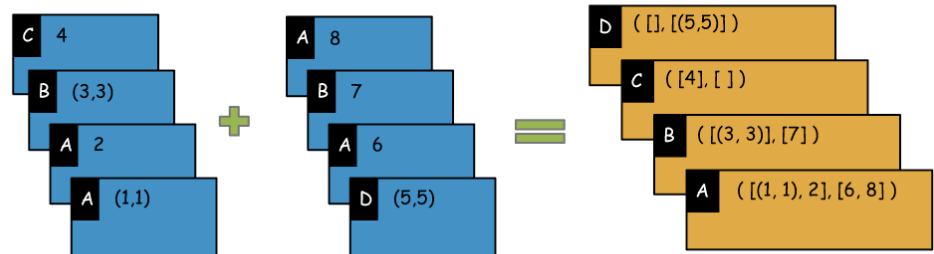
```
# ("A", [(1,1), 2], [6,8]))
```

```
# ("B", [(3,3)], [7])
```

```
# ("C", [4],[ ])
```

```
# ("D", [],[ (5,5)])
```

cogroup



Pair RDD actions

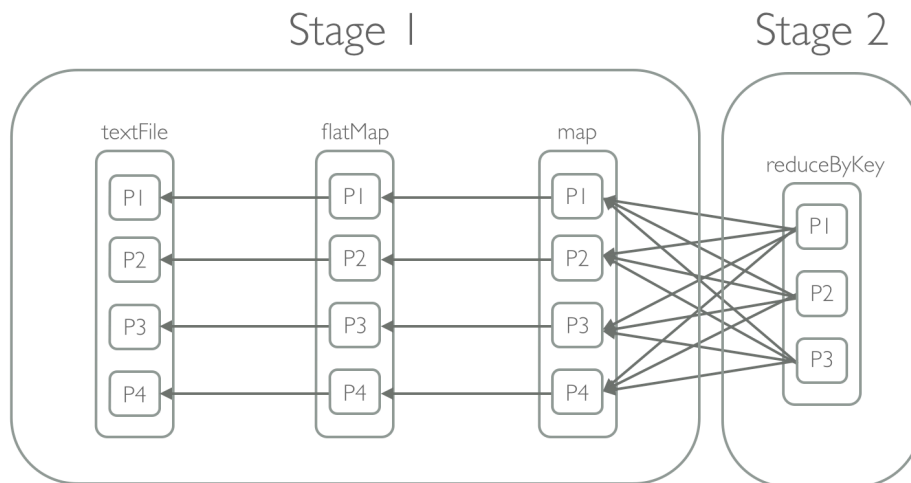
- Actions that make use of the key/value nature of pair RDDs

```
pets = sc.parallelize([("cat", 1), ("dog", 1), ("cat", 2)])  
pets.countByKey() # count the number of elements for each key  
# => {(cat, 2), (dog, 1)}
```

```
d = pets.collectAsMap() # Lookup through Python dictionary  
d["dog"]  
# => 1
```

Partitions

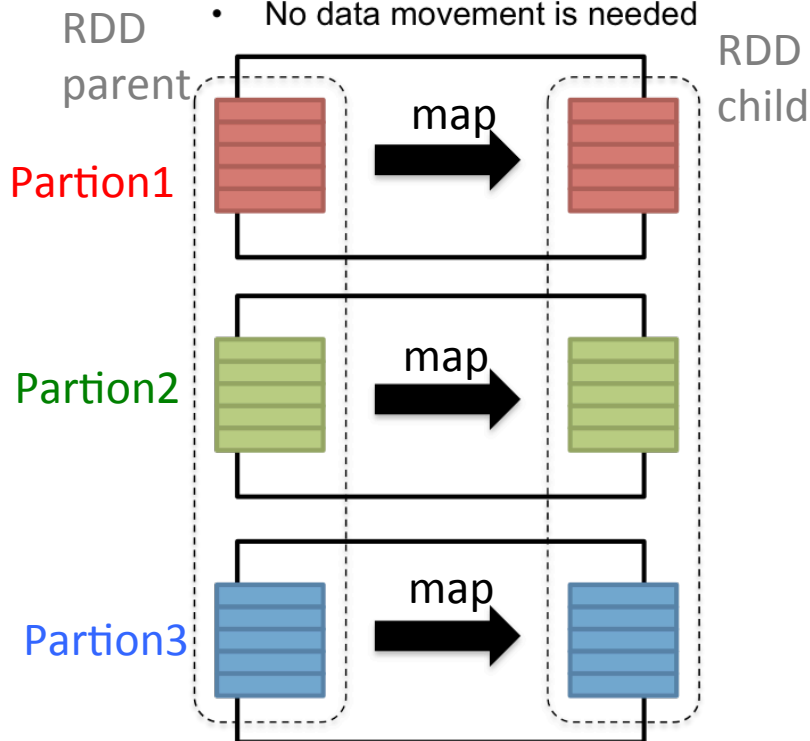
- Each RDD is split up into a number of partitions
- Parallelism is determined by the number of partitions
 - `rdd.getNumOfPartitions()`
- Why is important ?
 - Operations can be performed in parallel in each partition:
 - Textfile + flatMap + map operations can be performed in parallel in P1, P2, P3, P4
 - Operations that can run on the same partition are executed in stages



Narrow vs Wide transformations

Narrow transformation

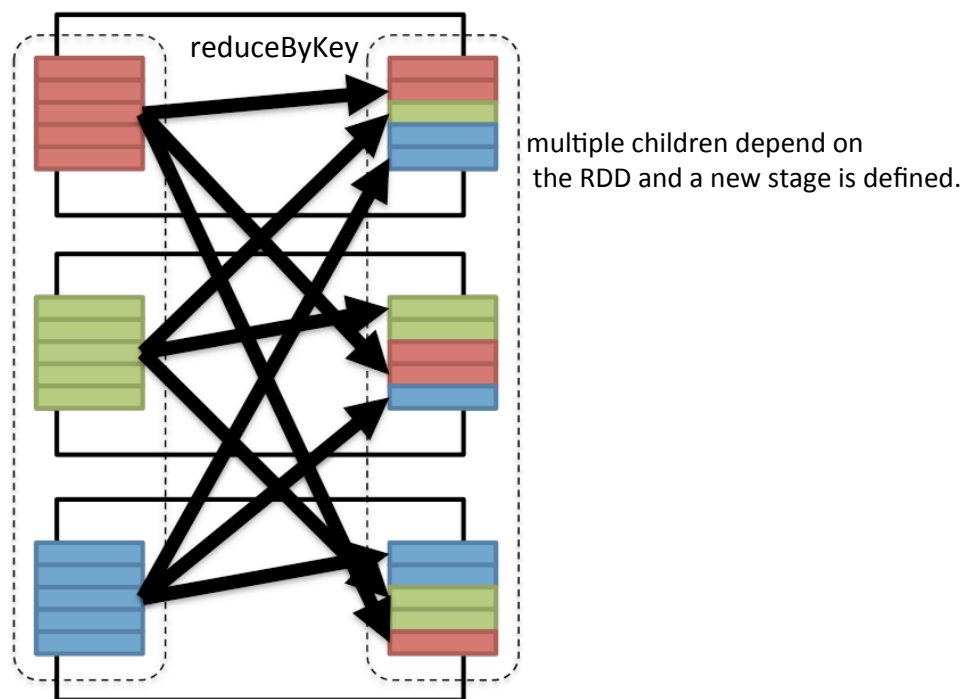
- Input and output stays in same partition
- No data movement is needed



e.g. `map`, `flatMap`, `filter`, `union`
each partition of the parent RDD is used
by at most 1 partition of the child RDD

Wide transformation

- Input from other partitions are required
- Data shuffling is needed before processing

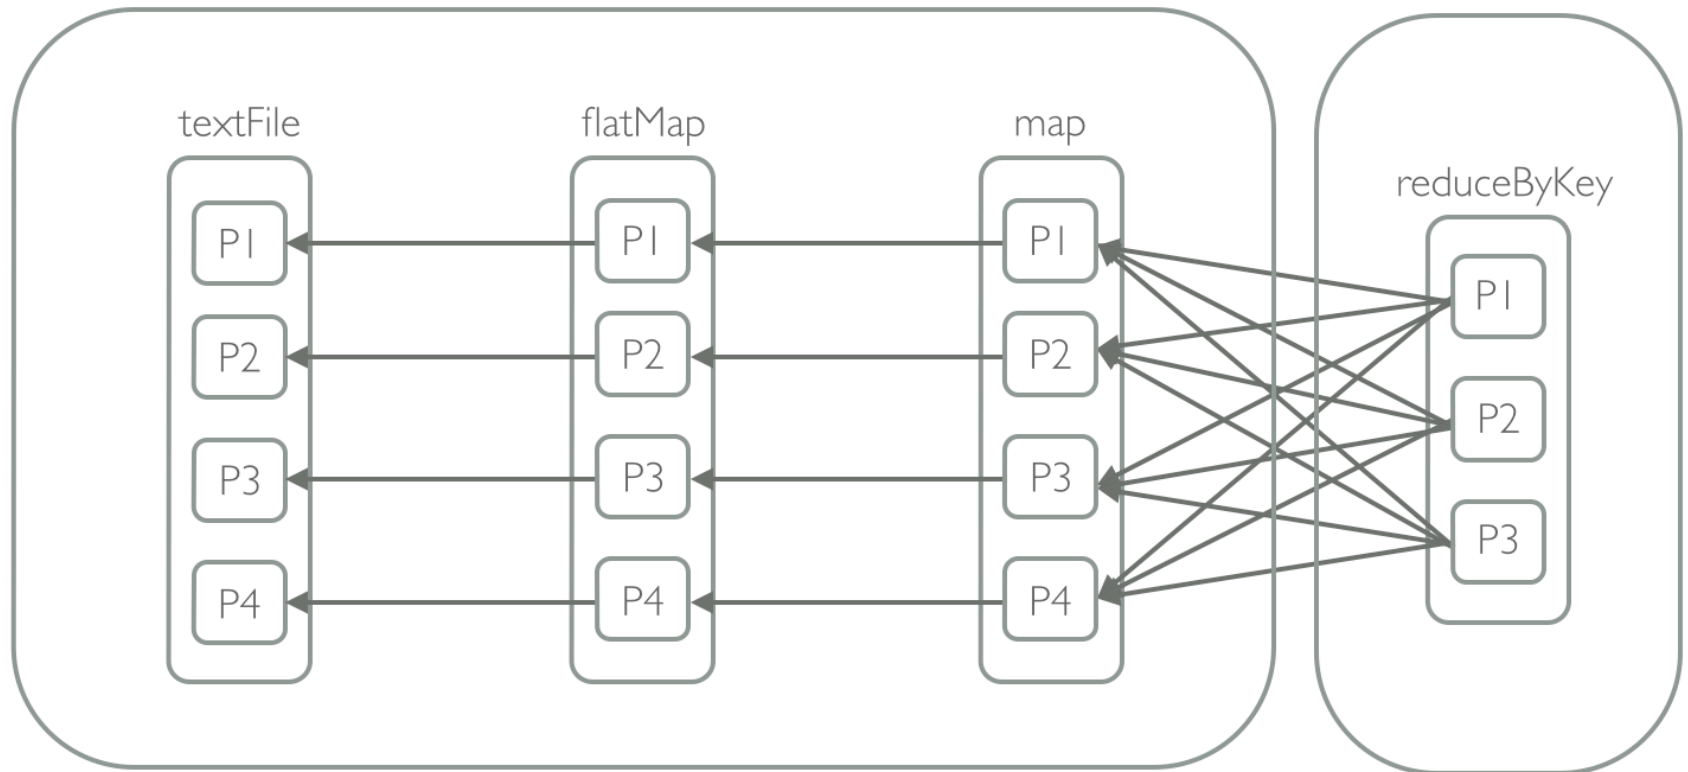


e.g. `groupByKey`, `distinct` or `join`
multiple child RDD partitions may depend
on a single parent RDD partition

Stages

Stage 1

Stage 2



Controlling the Level of Parallelism

- All the pair RDD operations take an optional second parameter for number of tasks

```
words.reduceByKey(lambda x, y: x + y, 5)
```

```
words.groupByKey(5)
```

```
visits.join(pageViews, 5)
```

Accumulators

- Accumulators are shared variables:
 - Used to aggregate values from workers nodes back to the driver node
 - Only the driver program can read the accumulator's value
 - Can be used as a counter or summations

```
gold = sc.accumulator(0)
```

```
def count_medals(events):  
    global gold, silver, bronze  
    for event in events:  
        if event.medal == 'GOLD':  
            gold.add(1)
```

```
results = sc.textFile('olympics.csv').filter(lambda x: x.country == 'USA')  
result.foreach(count_medals)
```

```
print gold.value
```


Broadcast Variables

- Spark's second type of shared variable
- Allows the program to efficiently send a large and read-only value to all the workers nodes.
 - Use it if your application needs to send a large, read-only lookup table or a large feature vector in a ML to all the nodes

```
import nltk
stopwords= set(nltk.corpus.stopwords.words('english'))
stopwords =sc.broadcast(stopwords)
if word in stopwords.value:
    pass
```

Example: PageRank

Why PageRank?

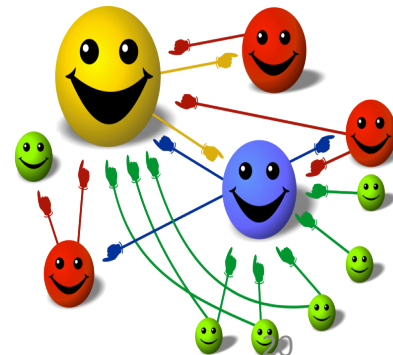
- Good example of a more complex algorithm
 - Multiple stages of map & reduce
- Benefits from Spark's in-memory caching
 - Multiple iterations over the same data

Basic Idea

- **PageRank** (PR) is an algorithm used by [Google Search](#) to rank websites in their search engine results.
- Give pages ranks (scores) based on links to them
 - Links from many pages → high rank
 - Link from a high-rank page → high rank

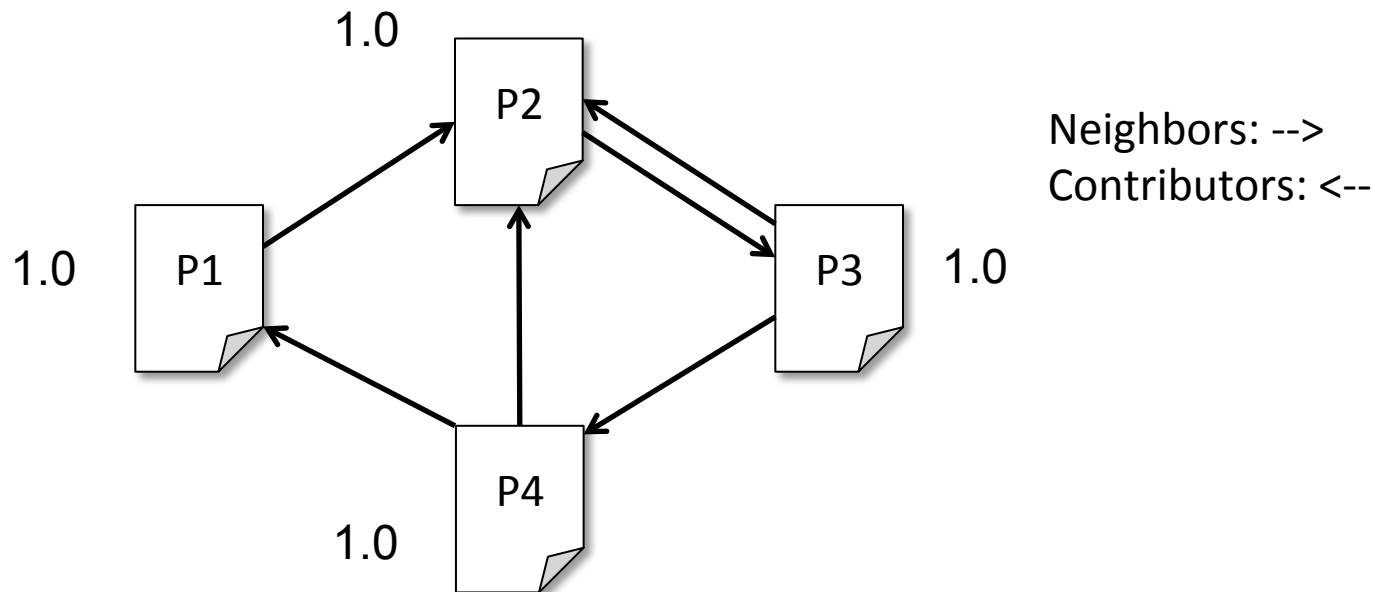
Overview of PageRank:

<https://en.wikipedia.org/wiki/PageRank>



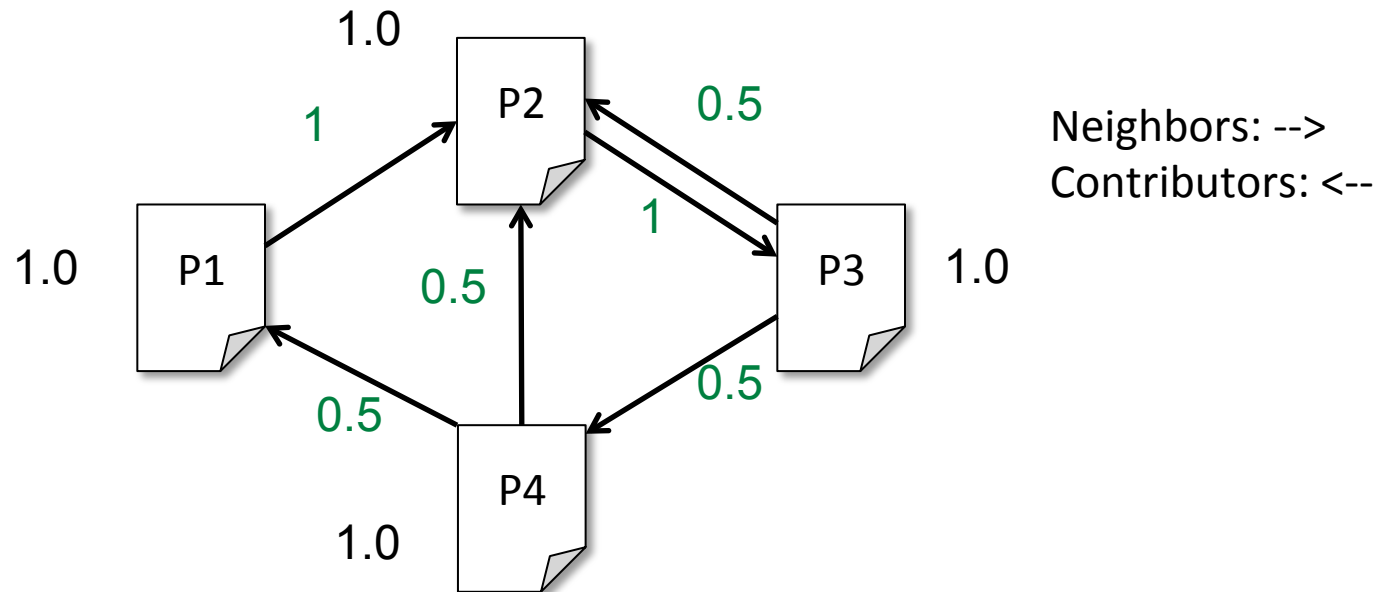
Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



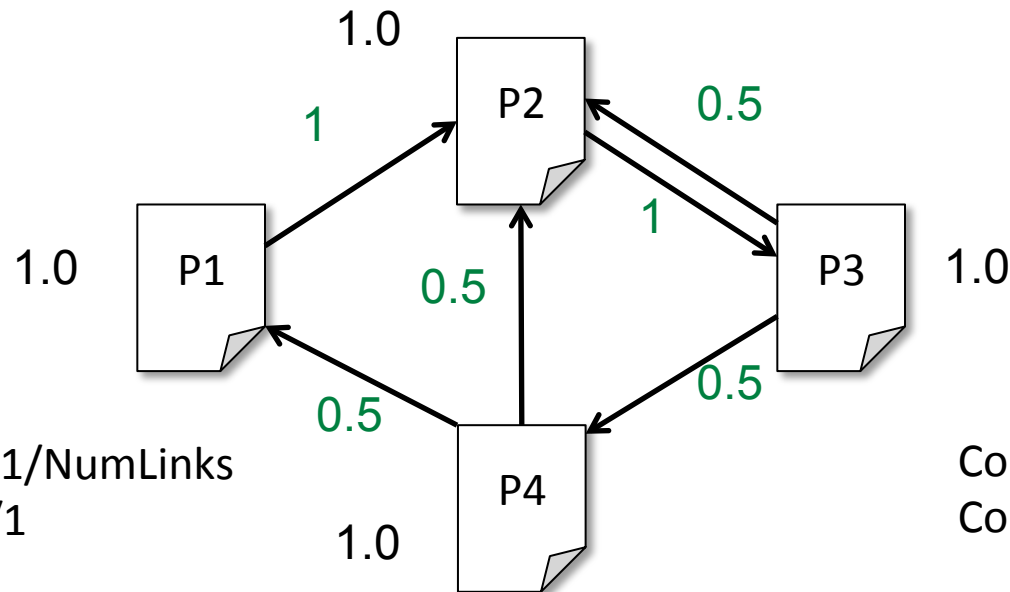
Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to “ \rightarrow ”).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$

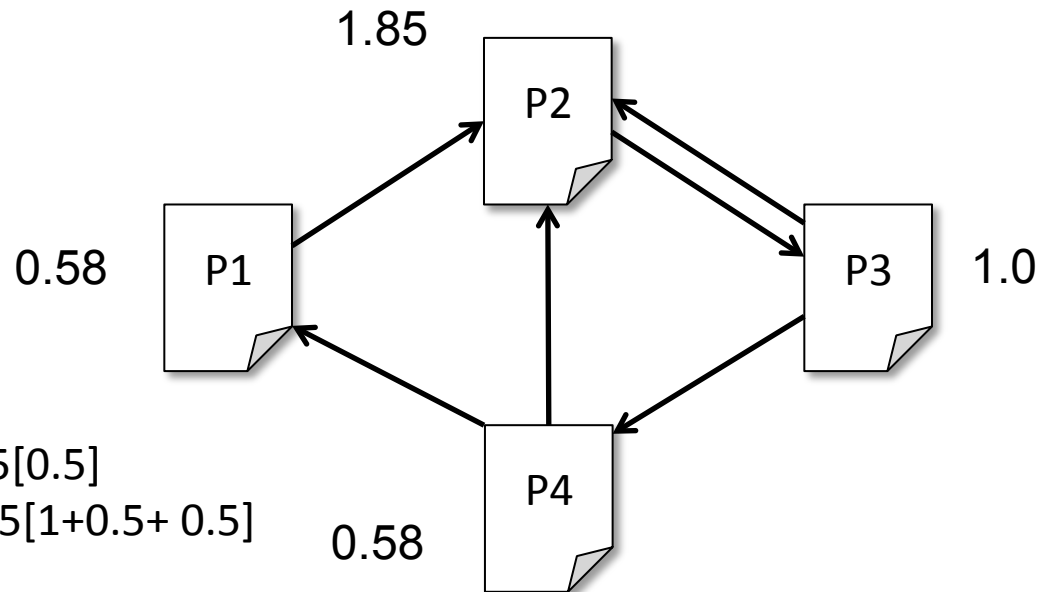


Contrib of P1: $PR1/\text{NumLinks}$
Contrib of P1 = $1/1$

Contrib of P4: $PR4/\text{NumLink}$
Contrib of P4 = $\frac{1}{2} = 0.5$

Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$

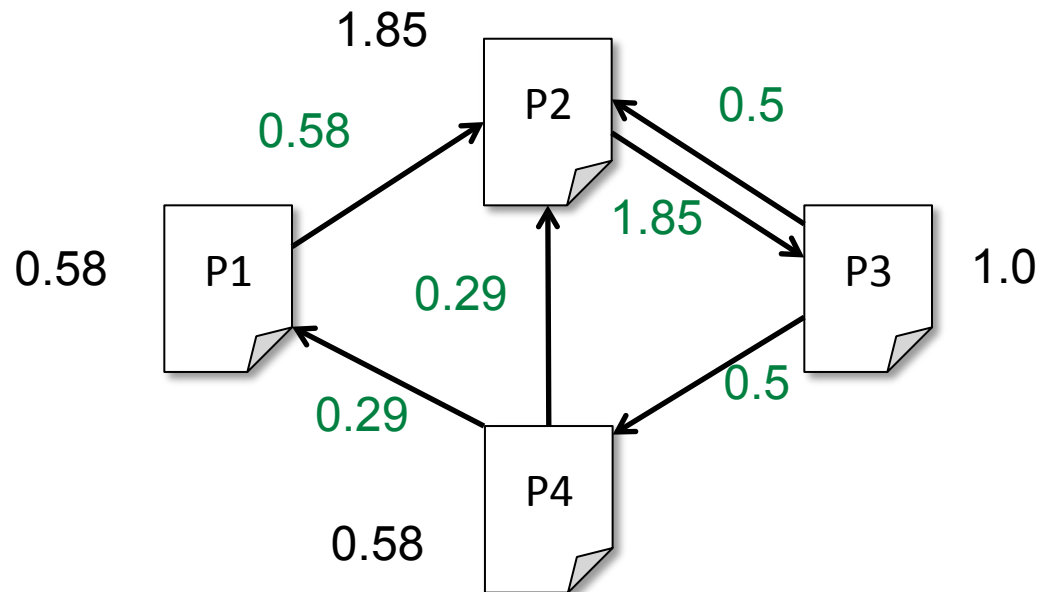


$$PR1 = 0.15 + 0.85[0.5]$$

$$PR2 = 0.15 + 0.85[1 + 0.5 + 0.5]$$

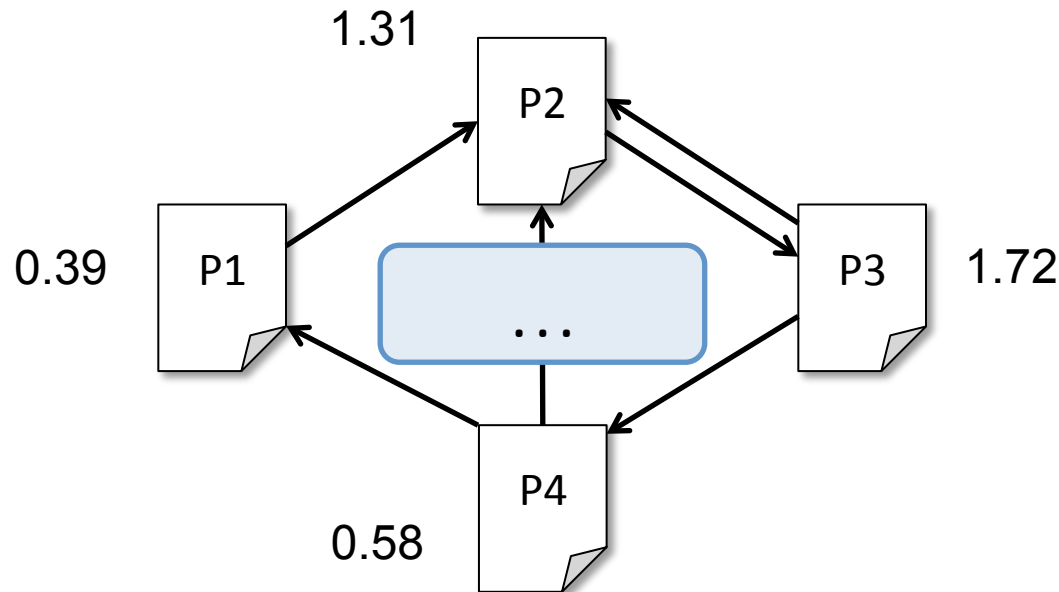
Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Algorithm

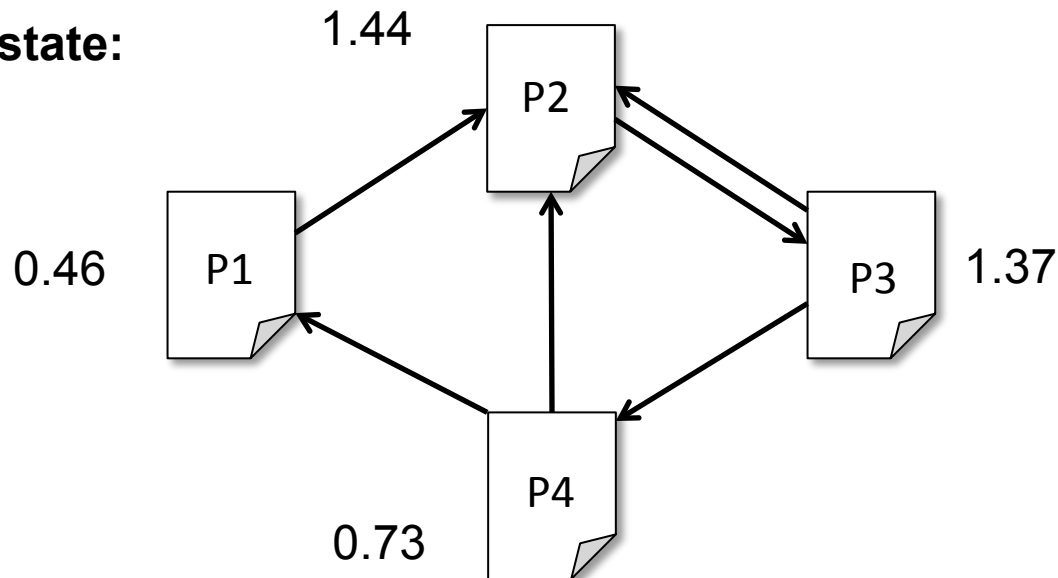
1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p send a contribution of $\text{rank}(p)/\text{numNeighbors}(p)$ to its neighbors (the pages it has links to).
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$

Final state:



PageRank Implementation

```
links = # RDD of (url, neighbors) pairs
ranks = # RDD of (url, rank) pairs

def compute_contribs(pair):
    [url, [links, rank]] = pair # split key-value pair
    return [(dest, rank/len(links)) for dest in links]

for i in range(NUM_ITERATIONS):
    contribs = links.join(ranks).flatMap(compute_contribs)
    ranks = contribs.reduceByKey(lambda x, y: x + y) \
        .mapValues(lambda x: 0.15 + 0.85 * x)

ranks.saveAsTextFile(...)
```

PageRank Implementation

```
links = sc.parallelize[("p1_url","p2"), ("p2_url","p3"), ("p3_url", ("p2","p4")),
("p4_url",("p1","p2"))
ranks = sc.parallelize[("p1_url",1), ("p2_url",1), ("p3_url",1), ("p4_url",1)]
```

```
def compute_contribs(pair):
    [url, [links, rank]] = pair # split key-value pair
    return [(dest, rank/len(links)) for dest in links]
```

```
for i in range(NUM_ITERATIONS):
    contribs = links.join(ranks).flatMap(compute_contribs)
    ranks = contribs.reduceByKey(lambda x, y: x + y) \
        .mapValues(lambda x: 0.15 + 0.85 * x)
```

```
ranks.saveAsTextFile(...)
```

join

[url, [links, rank]]

(dest, rank/len(links))

flatMap

("p1_url", ("p2",1)),
 ("p2_url", ("p3", 1),
 ("p3_url", (("p2","p4"), 1),
 ("p4_url", (("p1","p2"), 1))

["p1_url", ["p2"], 1],
 ["p2_url", ["p3"], 1],
 ["p3_url", ["p2","p4"], 1],
 ["p4_url", ["p1","p2"], 1]

P1 contribute to P2 → 1
 P2 contribute to P3 → 1
 P3 contribute to P2 and P4 → 0.5
 P4 contribute to P1 and P2 → 0.5

[(p2, 1), (p3,1), (p2, 0.5),
 (p4, 0.5), (p1 , 0.5), (p2, 0.5)]

reduceByKey

reduceByKey

[(p2, (1 +0.5 +0.5)), (p3,1), (p4, 0.5), (p1 , 0.5)]

[(p2,2), (p3,1), (p4, 0.5), (p1 , 0.5)]

mapValues

mapValues

[(p2, (0.15 + 0.85 *2)), (p3, (0.15 + 0.85 *1)), (p4, (0.15 + 0.85 *0.5)), (p1 , (0.15 + 0.85 *0.5))]

[(p2, 1.85), (p3, 1)), (p4, 0.58)), (p1 , 0.58)]

Logistic Regression

```
# Iterative machine learning algorithm
# Find best hyperplane that separates two sets of points in a
# multi-dimensional feature space. Applies MapReduce operation
# repeatedly to the same dataset, so it benefits greatly
# from caching the input in RAM

points = spark.textFile(...).map(parsePoint).cache()
w = numpy.random.randn(size = D) # current separating plane
for i in range(ITERATIONS):
    gradient = points.map(lambda p: (1 / (1 + exp(-
        p.y*(w.dot(p.x)))) - 1) * p.y * p.x).reduce(lambda a, b: a
        + b)
    w -= gradient
Print ("Final separating plane", w)
```

Best Practices

- Level of parallelism recommended: 3 tasks per CPU core.
- Reduce working set size
- Avoid groupByKey for associative operations
- Avoid reduceByKey when the input and output value types are different
- Avoid the flatMap-join-groupBy pattern
- Python memory overhead
- Use broadcast variables
- Cache judiciously
- Don't collect large RDDs
- Minimize amount of data shuffled
- Know the standard library
- Use dataframes

Complete article at:

<https://robertovitillo.com/2015/06/30/spark-best-practices/>