



PARTNERSHIP FOR ADVANCED COMPUTING IN EUROPE

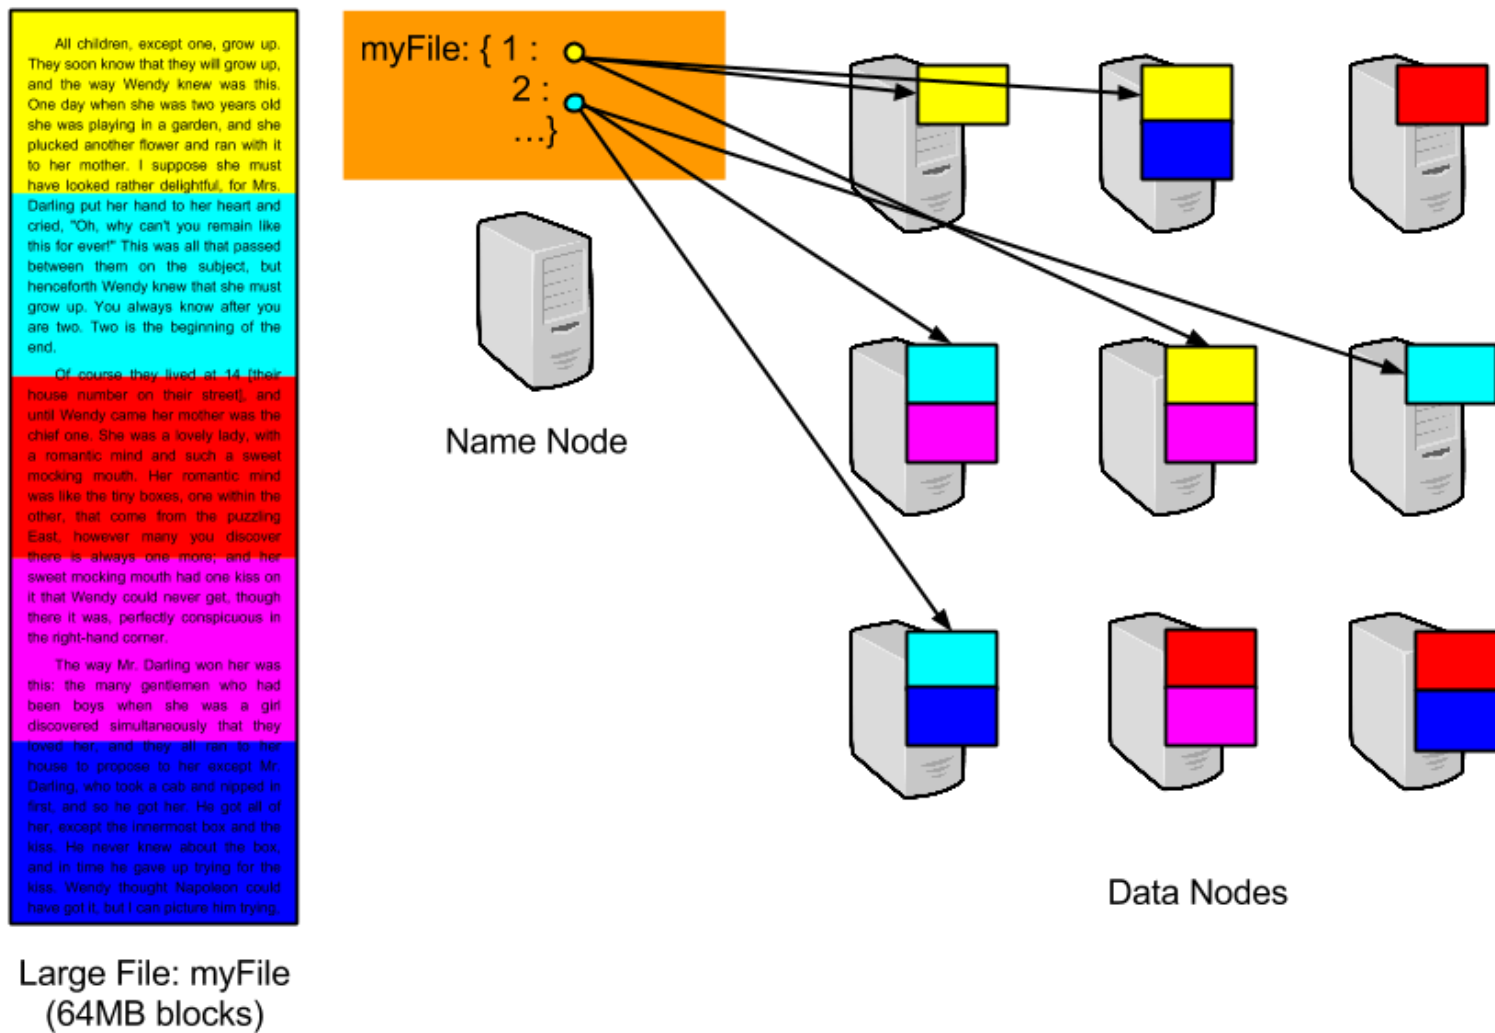
Hadoop

Amy Krause, Andreas Vroutsis

EPCC, The University of Edinburgh

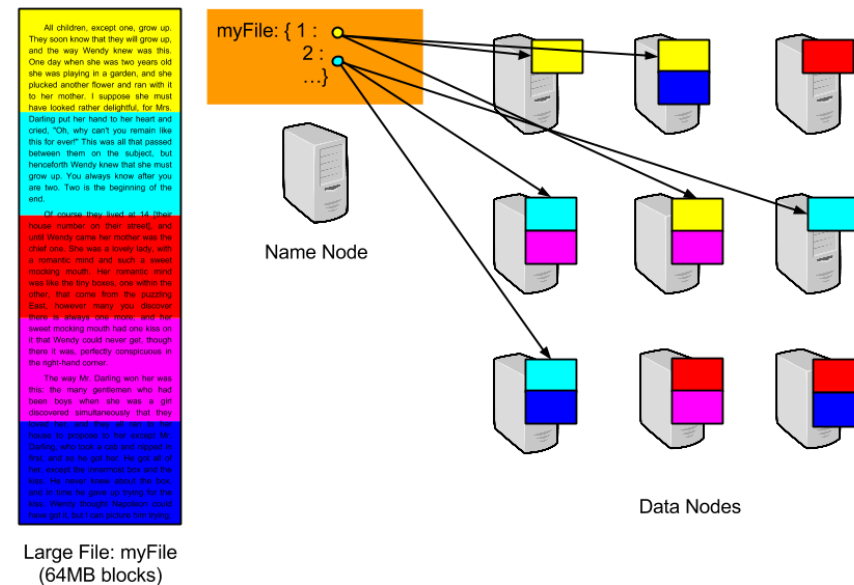
Slides thanks to Ally Hume, EPCC

Hadoop Distributed File System

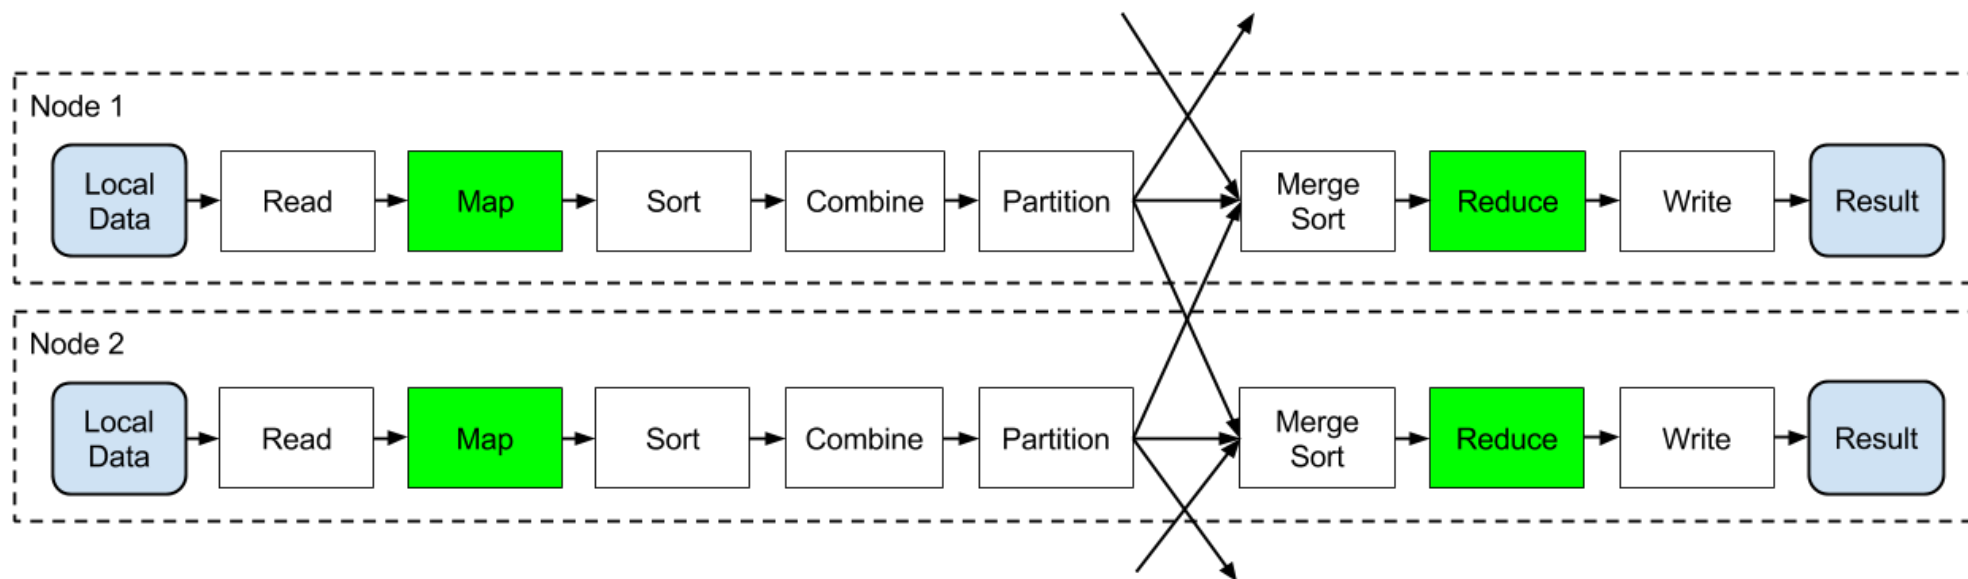


Hadoop Distributed File System

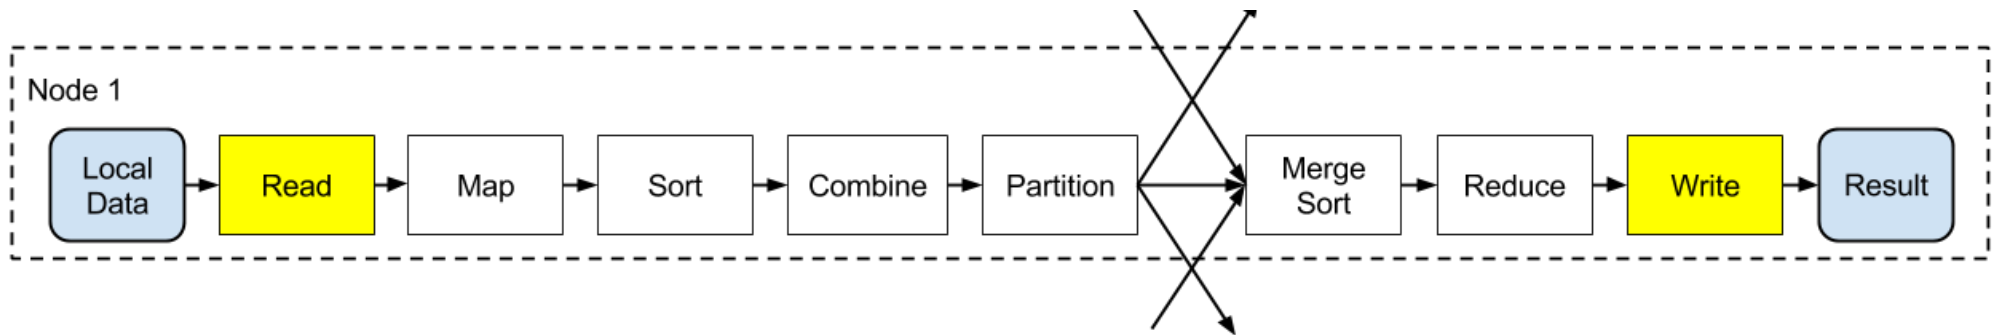
- ▶ Typical use: write once, read many
 - ▢ Computation runs on Data Nodes
- ▶ Distributed
- ▶ Data redundancy
- ▶ Cluster of commodity nodes
- ▶ Designed to withstand failure
 - ▢ But Name Node is a single point of failure (see secondary name node)
- ▶ Not a POSIX file system
- ▶ Placement strategies can be aware of data centre configuration



Hadoop Framework



Reading and writing the data



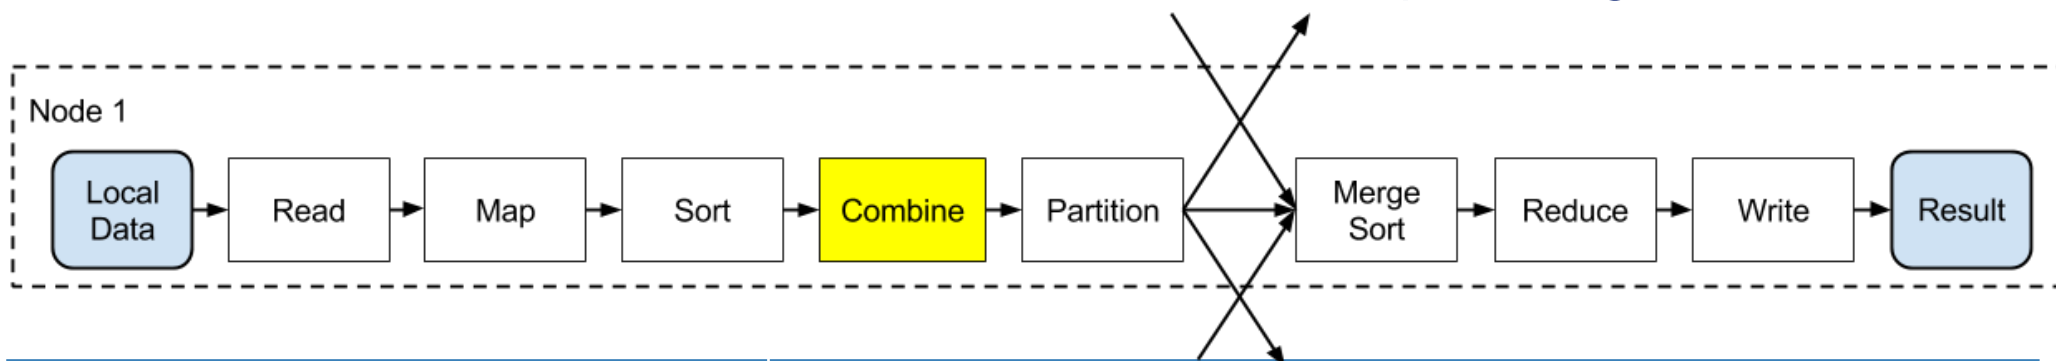
▶ InputFormat interface

- ▶ TextInputFormat (key: byte offset of line, value: line text)
- ▶ KeyValueTextInputFormat (each line has key/seperator/value)
- ▶ SequenceFileInputFormat (Hadoop's compressed binary format)
- ▶ NLineInputFormat (like TextInputFormat but multi-line)

▶ OutputFormat interface

- ▶ TextOutputFormat (one record per line, key/seperator/value)
- ▶ SequenceFileOutputFormat (compressed binary)
- ▶ Filename is "part-xxxx" where xxxx is the partition ID

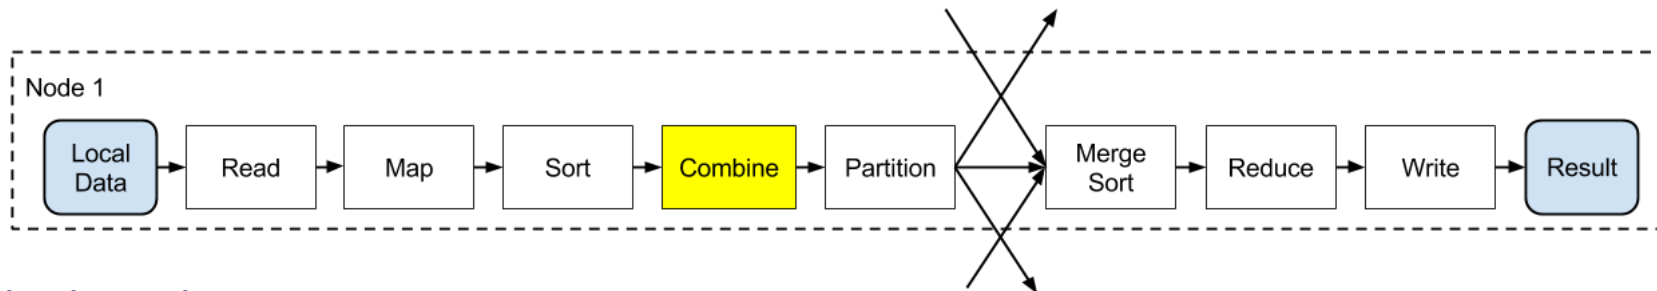
Optimising with a combiner



Map Input	Map Output
<0 : "A boy drove a car">	[<a,1>, <boy,1>, <drove,1>, <a,1>, <car,1>]
<1 : "A car drove at a bus">	[<a,1>, <car,1>, <drove,1>, <at,1>, <a,1>, <bus,1>]
<2 : "Can a boy drive a car?">	[<can,1>, <a,1>, <boy,1>, <drive,1>, <a,1>, <car,1>]
<3 : "A danger – a banana!">	[<a,1>, <danger,1>, <a,1>, <banana,1>]

Combiner Input	Combiner output
<a,[1,1]>	<a, [2]>
<boy, [1,1]>	<boy, [2]>
<car,[1,1,1]>	<car, [3]>
<drove,[1,1]>	<drove, [2]>

Combiner properties




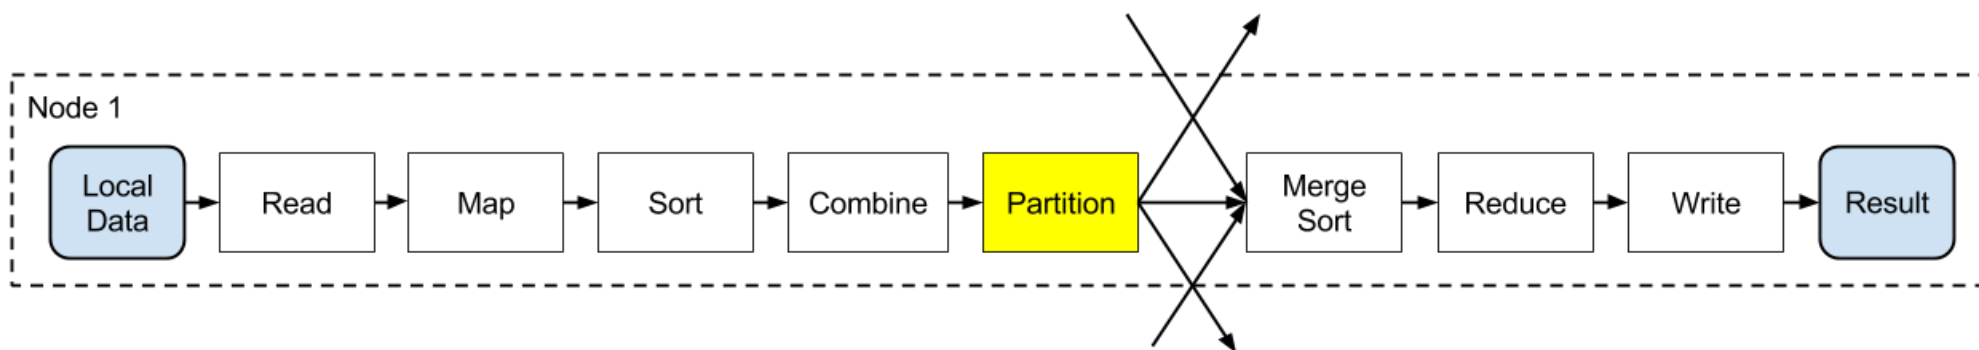
- ▶ Optimisation only
 - ▶ Framework may execute zero, one or more times
 - ▶ Must not alter the final result
 - ▶ A helper to the reducer
- ▶ Keys *must* not be altered
 - ▶ Hadoop does not re-sort after the Combine stage

	Input	Output
Map	<Key1 : Value1>	List(<Key2 : Value2>)
Combine	<Key2 : List(Value2) >	<Key2 : List(Value2) >
Reduce	<Key2 : List(Value2) >	List(<Key3 : Value3>)



Partitioner


- ▶ Hash Partitioner
 - ▶ Default
- ▶ Total Order Partitioner
 - ▶ Maintains order
 - ▶ Configure to partition evenly
- ▶ Bespoke
 - ▶ For highly skewed data hash partitioner may not partition work evenly
 - ▶ Maybe some keys require more processing by Reducer





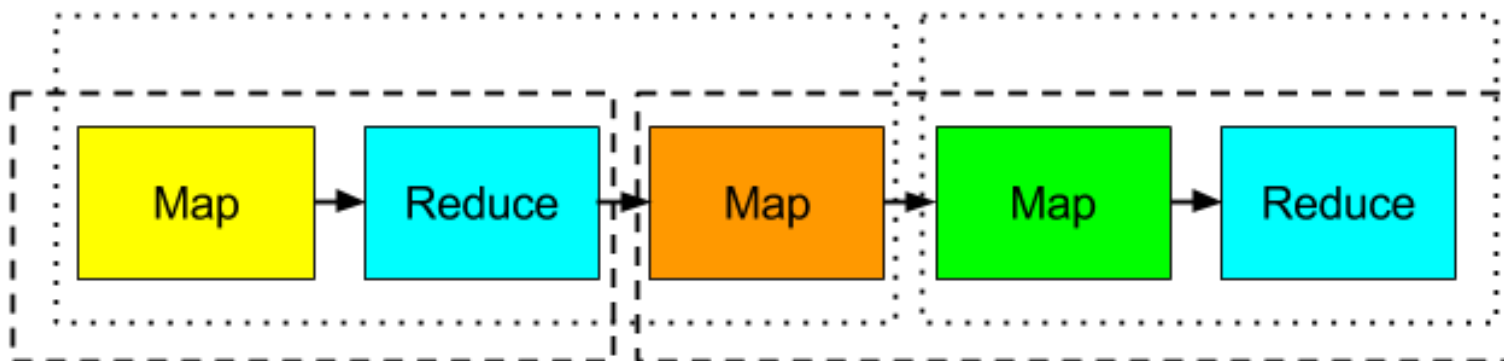
Chaining MapReduce Jobs

- ▶ A single map reduce job has
 - ▶ One REDUCE stage
 - ▶ One or more MAP stages before the reduce
 - ▶ Zero or more MAP stages after the reduce



Chaining MapReduce Jobs

- ▣ Need to chain multiple map reduce jobs when:
 - ▣ There is more than one REDUCE stage (grouping of data by key)
 - ▣ MAP stages between REDUCE jobs could be part of either job





Chain, but don't iterate



- ▶ Each Hadoop job reads data from the HDFS and writes output to the HDFS

- ▣ No data is maintained in memory between jobs

- ▶ Fine for short chains of processing

- ▶ Very inefficient for iterative algorithms

- ▣ Data (even static data) must be read from disk at each iteration

- ▶ Spark – supports caching data

- ▶ Twister – iterative map reduce





Programming Hadoop

- ▶ Hadoop framework is written in Java
- ▶ Two models for writing Map, Reduce and Combine functions
 - Java classes
 - Hadoop streaming
 - Functions are scripts that read from standard input and write to standard output
- ▶ If writing your own partitioners or getting into the internals of Hadoop you will need to use Java
 - But for most problems you do not need to do this.

Map class in Java

Must implement function:
void map(InputKeyType, InputValueType, Context)

**Mapper<InputKeyType, InputValueType,
OutputKeyType, OutputValueType >**

```
public static class MapClass
    extends Mapper<Text, Text, Text, Text>
{
    public void map(Text key, Text value, Context context)
    {
        context.write(value, key);
    }
}
```

**This mapper simply swaps
the key and value**

**write output data using context.write(outputKey,
outputValue)**

**Can call multiple times and hence output
List(<OutputKeyType, OutputValueType>)**



Reduce class in Java

```
public static class Reduce
    extends Reducer<Text, Text, Text, Text>
{
    public void reduce( Text key,
                       Iterable<Text> values,
                       Context context)
    {
        String csv = "";
        for (Text val:values)
        {
            if (csv.length() > 0) csv += ",";
            csv += val.toString();
        }
        context.write(key, new Text(csv));
    }
}
```


**Reducer<
InputKeyType, InputValueType,
OutputKeyType, OutputValueType >**

**Uses iterator to get list of
values – can thus support
large lists with low memory
footprint. So long as the rest
of the method is similarly low
memory. This example is not!**

**write output data using context.write(outputKey,
outputValue)
Can call multiple times if desired**

Streaming Mapper

```
1<TAB>A long time ago
2<TAB>in a galaxy far
3<TAB>far away
```



```
a<TAB>1
long<TAB>1
time<TAB>1
ago<TAB>1
in<TAB>1
a<TAB>1
galaxy<TAB>1
...
```

- ▶ Input: rows of key/value pairs separated by TAB character
- ▶ Output: rows of key/value pairs separated by TAB character
- ▶ **Stateless**
 - ▢ Process one line at a time with no state maintained between lines.




Streaming Reducer

- ▶ Input is rows of key/value pairs separated by TAB character
- ▶ Input guarantees that all the key/value pairs associated with a specific key will be contiguous in the input stream
 - ▶ When key changes you know you have seen all the values associated with that key
- ▶ Output rows of key/value pairs separated by TAB character
- ▶ **Stateless**
 - ▶ Can maintain state while processing rows with the same key.
 - ▶ Must not maintain state across rows with different keys

Streaming Reducer

```
a<TAB>1  
a<TAB>1  
a<TAB>1  
far<TAB>1  
far<TAB>1  
time<TAB>1
```



```
a<TAB>3  
far<TAB>2  
time<TAB>1
```




Hadoop vs MPI/HPC

- ▶ Fault tolerance
 - ▶ Hadoop is designed specifically with fault tolerance in mind
 - ▶ MPI provides little support for fault tolerance and most MPI programs assume the system hardware will not fail
- ▶ Specific vs general
 - ▶ Hadoop is a framework for a specific data processing pattern
 - ▶ MPI allows you to code any algorithm you wish
- ▶ Iterative algorithms
 - ▶ Hadoop very poor at multiple iterations over the data
 - ▶ Very easy to write such programs in MPI
- ▶ Speed
 - ▶ If you have a reliable HPC system an optimised MPI implementation should perform considerably better than a Hadoop solution



Hadoop vs MPI/HPC cont.

- ▶ Hadoop good when data written once – processed often
- ▶ Hadoop uses key value pair structures that can be fragmented in memory leading to poor cache efficiency
- ▶ Trade off between simple, highly scalable on commodity hardware against highly optimised implementation on very expensive hardware



Hadoop vs MPI/HPC cont.

▣ Cost

- ▣ Hadoop simple to write and can run reliably on commodity hardware.
- ▣ MPI typically run on expensive HPC systems
 - ▣ MPI can run on clouds but have to build your own fault tolerance.

▣ Dynamic nature of data

- ▣ Hadoop is good for processing massive amounts of data that is written once and processed often
- ▣ HPC systems may not scale well to such massive datasets being uploaded.

APACHE HBASE

Hadoop Ecosystem

▶ HBASE

- ▶ Distributed, scalable big data store
- ▶ Columnar database

▶ PIG

- ▶ Higher level data flow language for programming Hadoop

▶ Mahout

- ▶ Scalable machine learning and data mining over Hadoop

▶ Spark



- ▶ Machine learning algorithms





A little more on Spark

- ▶ Explicitly supports caching data
 - Speeds up iterative algorithms
- ▶ Can use HDFS as the data source
- ▶ More than just map/reduce
 - Transformations:
 - map, filter, union, Cartesian, join, sample...
 - Actions:
 - reduce, collect, count, first, countBy, foreach...



Additional reading

▣ Google File System

- ▣ <http://static.googleusercontent.com/media/research.google.com/en/archive/gfs-sosp2003.pdf>

▣ Map Reduce

- ▣ <http://static.googleusercontent.com/media/research.google.com/en/archive/mapreduce-osdi04.pdf>

▣ Examples taken from Hadoop in Action

- ▣ <http://www.manning.com/lam/>

▣ For Hadoop 3 I like O'Reilly's Hadoop, The Definitive Guide

▣ Plenty online – course does not require material outside of the lectures and practical.

