



PARTNERSHIP FOR ADVANCED COMPUTING IN EUROPE

Spark Cluster Overview


Amy Krause, Andreas Vroutsis

EPCC, The University of Edinburgh

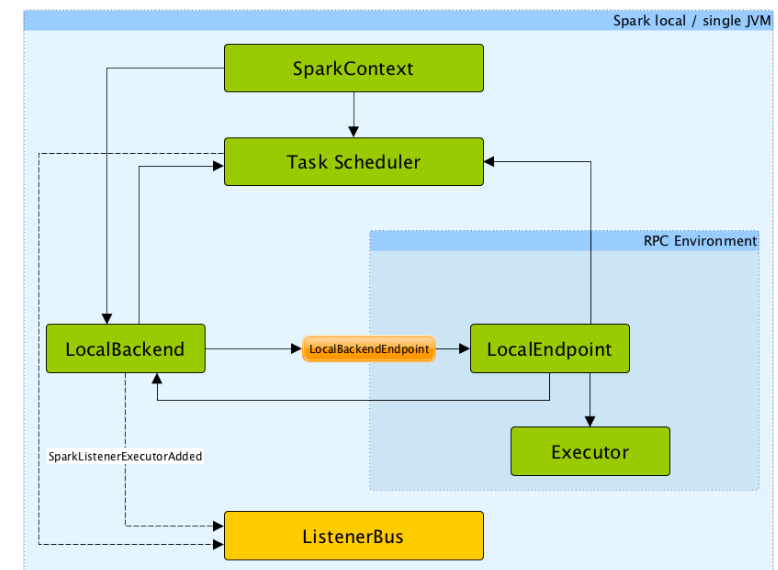
Spark Execution modes

It is possible to run a Spark application using cluster mode, local mode (pseudo-cluster) or with an interactive shell (pyspark or spark-shell).

Cluster mode

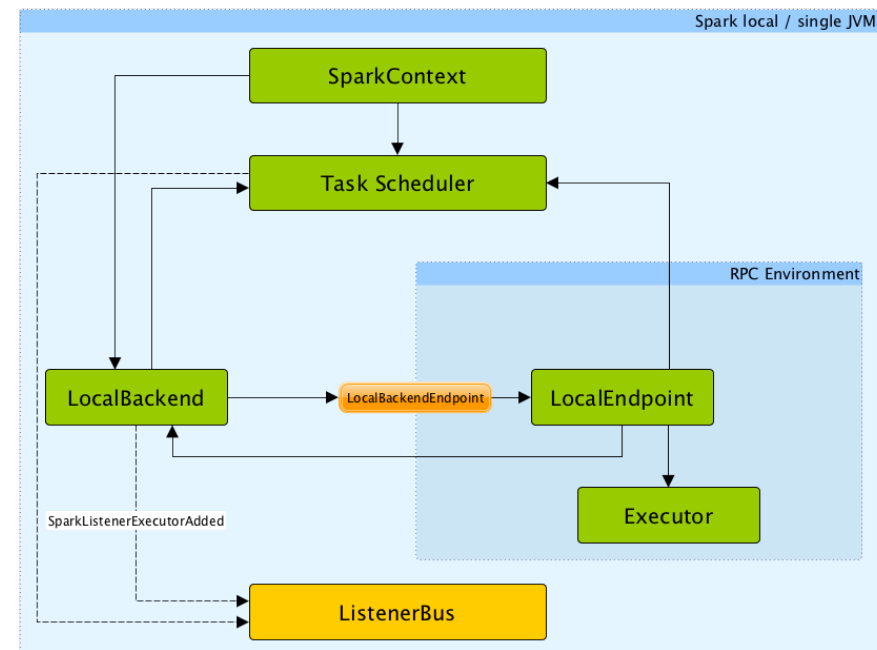


Local mode




Spark Execution – Local Mode

- ▶ In this non-distributed single-JVM deployment mode.
- ▶ Spark spawns all the execution components - driver, executor, LocalSchedulerBackend, and master - in the same single JVM.
- ▶ The default parallelism is the number of threads as specified in the master URL.




Standalone Deploy Mode

- ▶ Simplest way to deploy Spark on a private cluster
 - ▶ Apache Mesos
 - ▶ Hadoop YARN
 - ▶ Kubernetes



- ▶ Spark is agnostic to the underlying cluster manager

Spark Execution – Cluster mode




- ▶ Spark applications are run as independent sets of processes, coordinated by a SparkContext in a *driver* (*) program.
- ▶ The *context* connects to the cluster manager *which allocates resources*.
- ▶ Each *worker* in the cluster is managed by an *executor*.
- ▶ The *executor* manages computation as well as storage and caching on each machine.

(*) driver → process running the main() function of the application and creating the SparkContext

Spark Execution – Cluster mode


- ▶ The application code is sent from the *driver* to the *executors*, and the executors specify the context and the various *tasks* to be run.
- ▶ The *driver* program must listen for and accept incoming connections from its executors throughout its lifetime.



Spark App

Each SparkContext creates a Spark application, which includes a lot of scheduling components.

Upon an **Action**, the driver program submits the job to the cluster manager.



Cluster manager

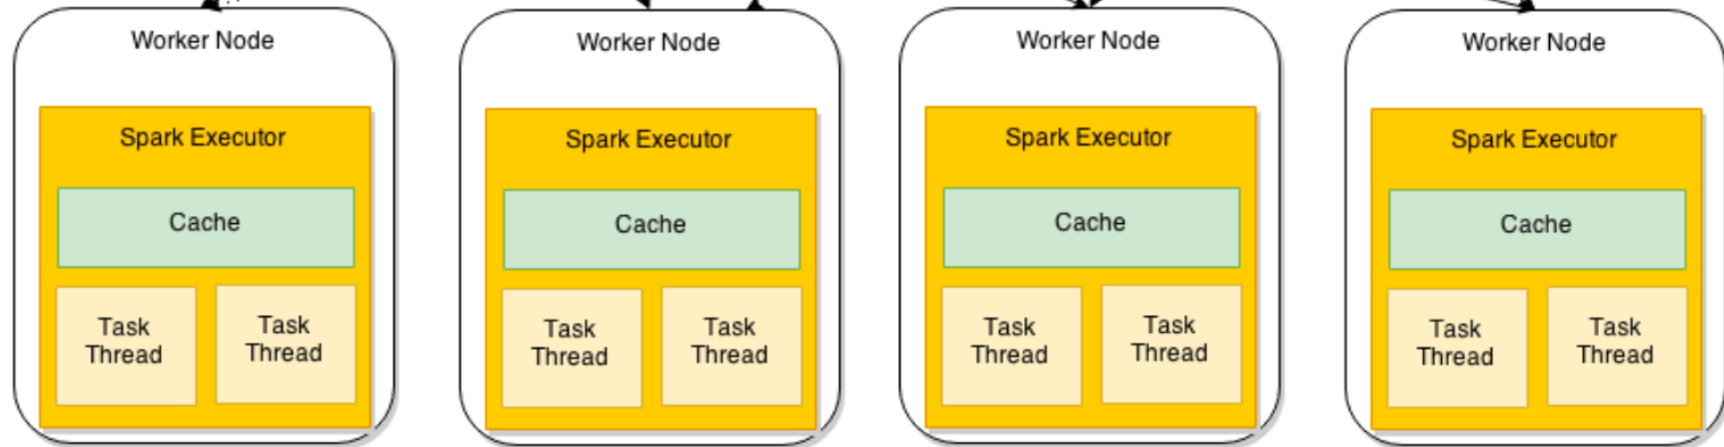
Start executors on Worker Nodes.

It does **not** know about stages.

Worker

Launch Spark Executor in a process.

Tasks are launched in separate threads, one per each core on the worker node (can be configured)



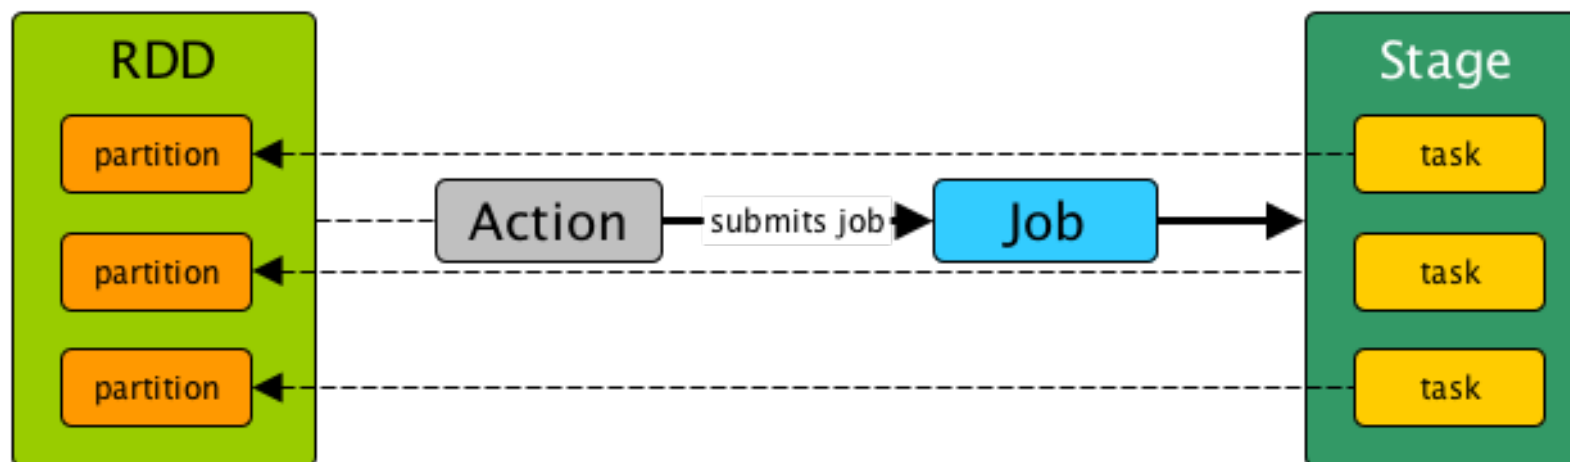


Spark: Standalone cluster – deploy modes

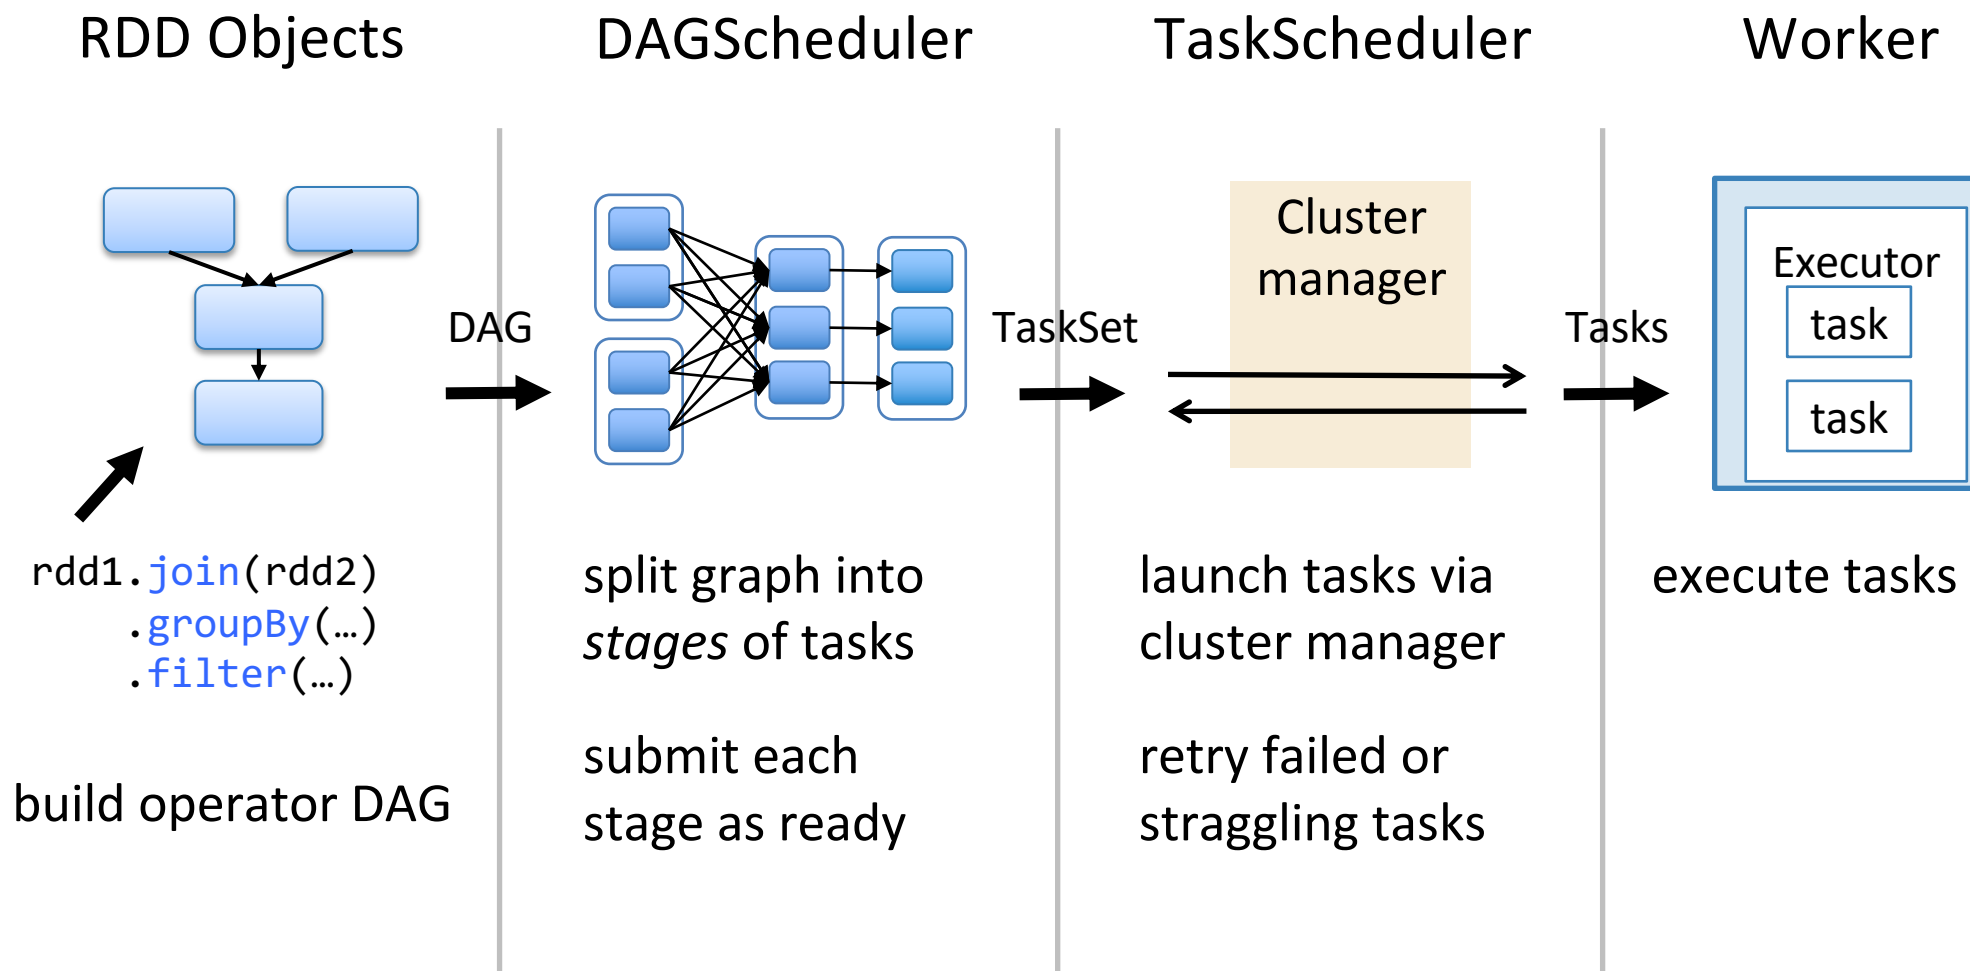
- ▶ For standalone clusters supports two deploy modes. They distinguish where the driver process runs:
 - ▶ *Client mode* (by default): the driver is launched in the same process as the client that submits the application.
 - ▶ *Cluster mode*: the driver is launched from one of the Worker processes inside the cluster.
 - ▶ The client process exits as soon as it fulfils its responsibility of submitting the application without waiting for the application to finish.
- ▶ Note: Currently, the **standalone mode** does not support **cluster mode** for **Python applications**.

Spark Components

- ▶ Task: individual unit of work sent to one executor over a sequences of partitions
- ▶ Job : set of tasks executed as a result of an action
- ▶ Stage: set of tasks in a job that can be executed in parallel – at partition level
- ▶ RDD: Parallel dataset with partitions
- ▶ DAG: Logical Graph of RDD operations



Job scheduling





Spark Application – wordcount.py

- ▶ The application that we are going to create is a simple “wordcount”:
 - ▶ Performs a **textFile** operation to read an input file in HDFS
 - ▶ **flatMap** operation to split each line into words
 - ▶ **map** operation to form (word, 1) pairs
 - ▶ **reduceByKey** operation to sum the counts (all the '1') for each word




Spark Application – wordcount.py

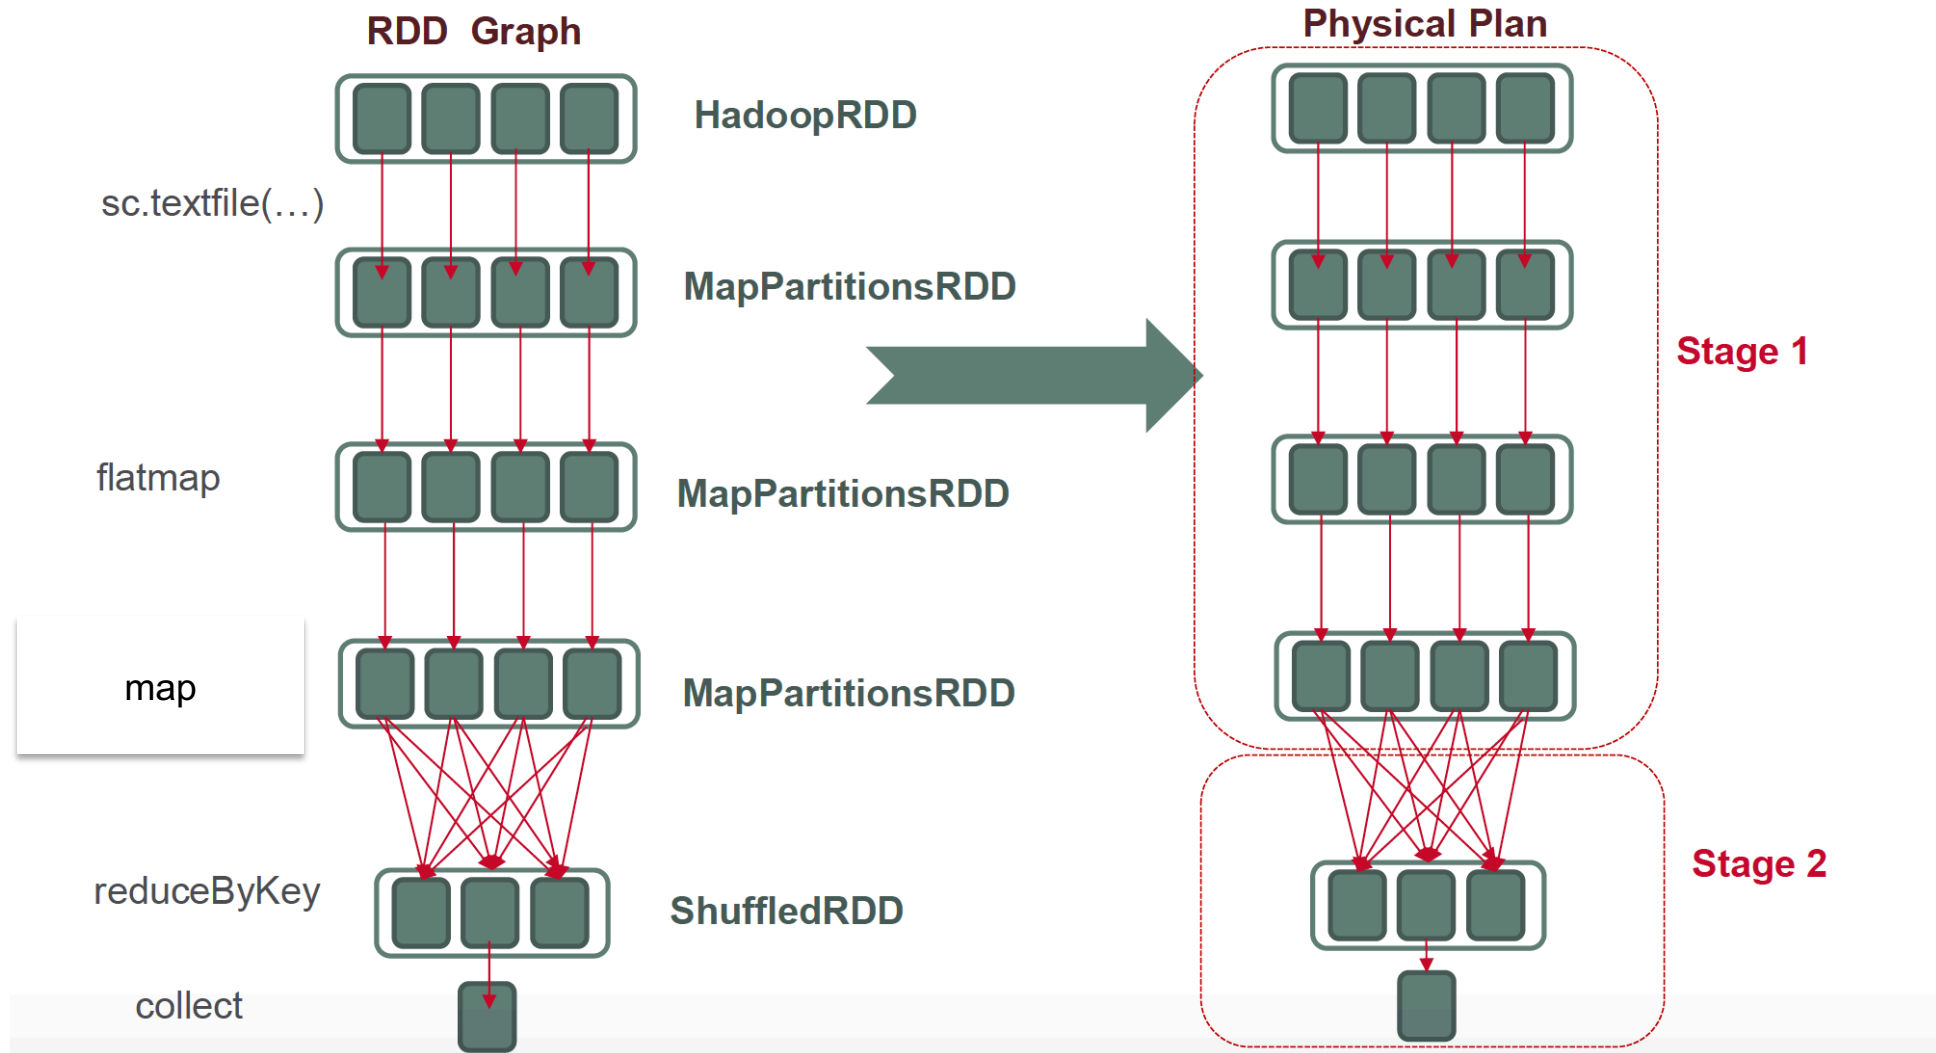
```
import sys
from pyspark import SparkContext, SparkConf

if __name__ == "__main__":
    conf = SparkConf().setAppName("Spark Count")
    sc = SparkContext(conf=conf)
    inputFile = sys.argv[1]
    textFile = sc.textFile(inputFile)
    wordCounts = textFile.flatMap(lambda line: line.split()).\
        map(lambda word: (word, 1)).reduceByKey(lambda a, b: a+b)
    output=wordCounts.collect()
    for (word, count) in output:
        print("%s: %i" % (word, count))
```

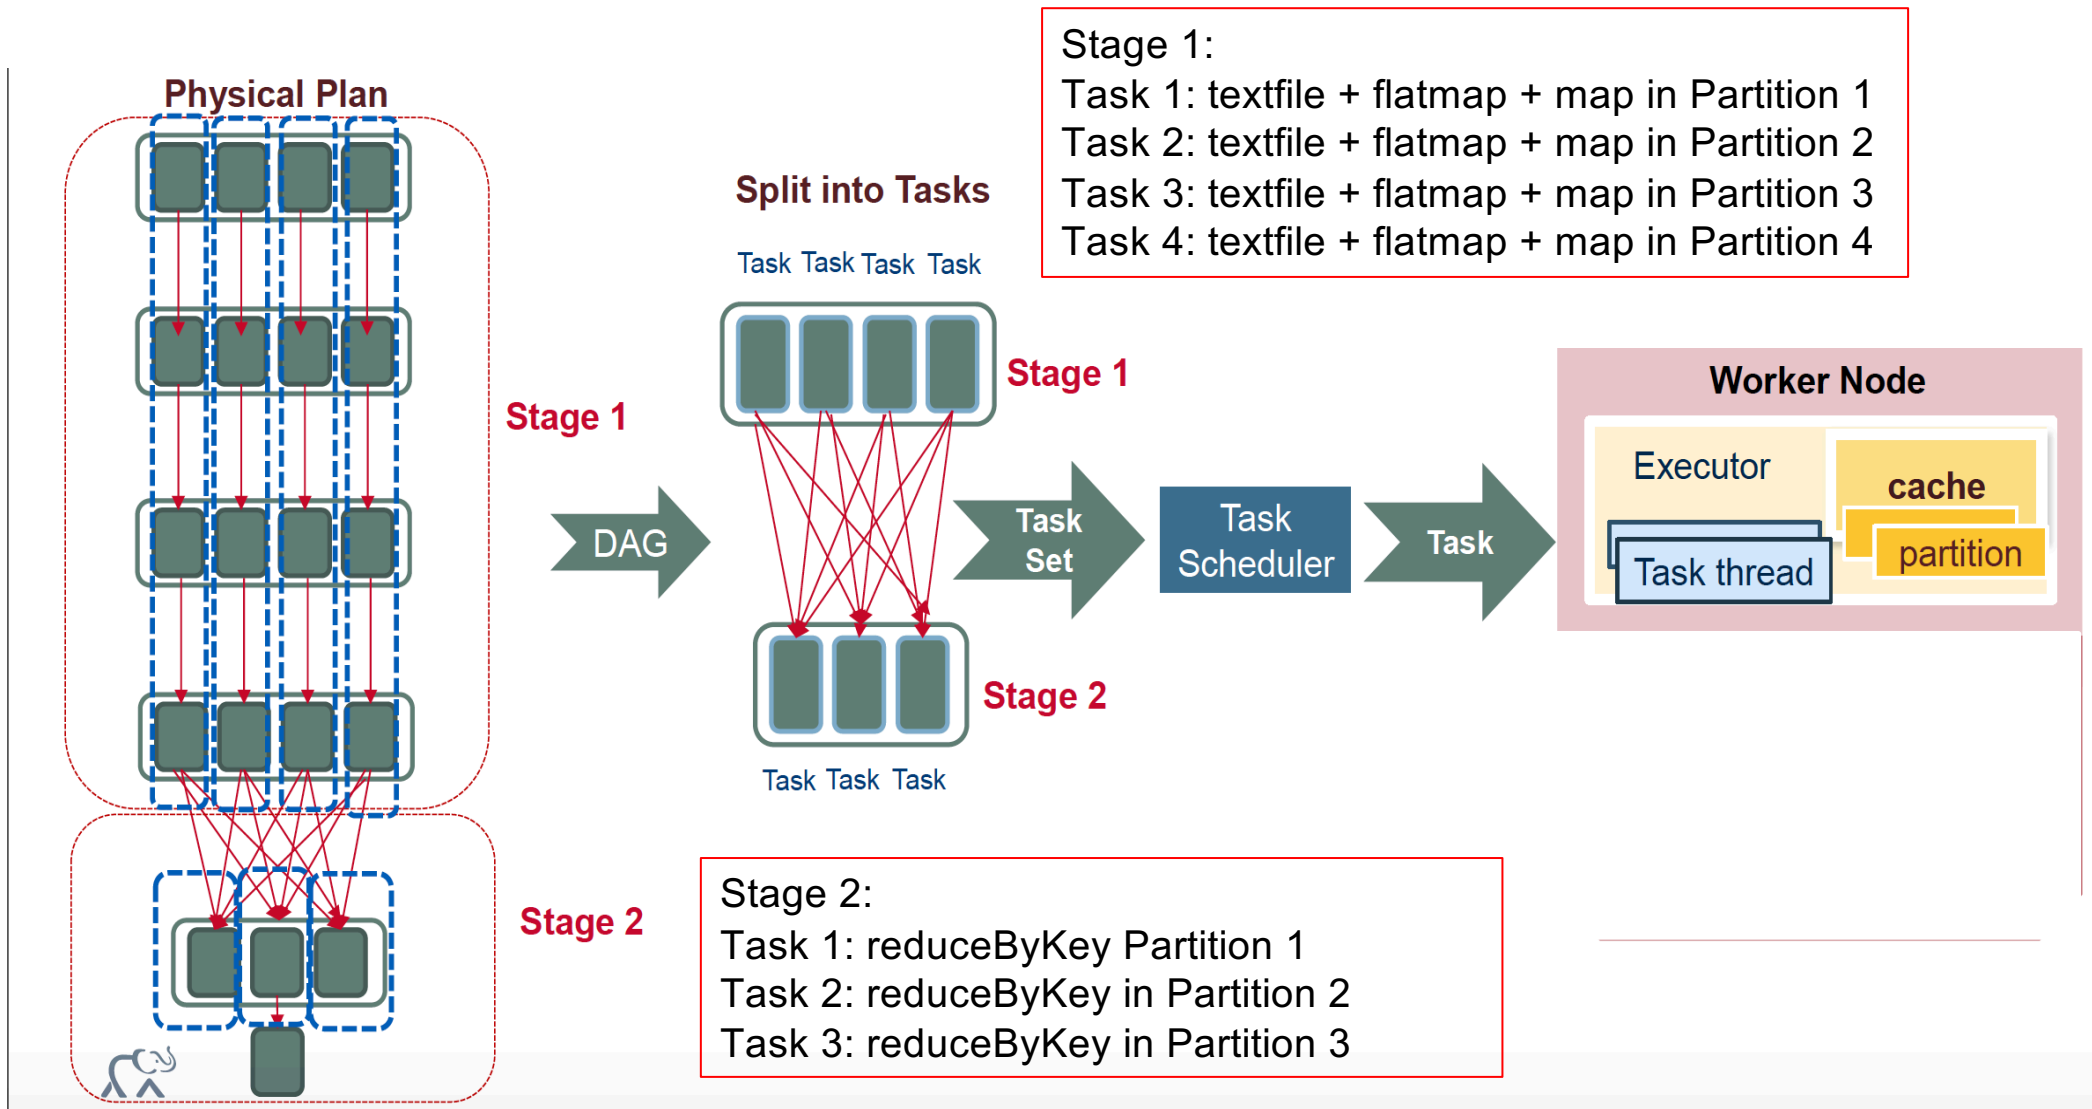
Spark Application – wordcount.py



RDD DAG -> Physical Execution plan



Initial RDD distributed among 4 partitions. Final RDD distributed among 3 partitions



Operations that can run on the same partition are executed in stages



Running Spark Applications

- ▶ **Notebooks** are great for:
 - ▶ developing and testing quickly experiment with the data
 - ▶ demos and collaborating with other people
- ▶ **Spark-submit** jobs are more likely to be used in **production**.



Running Spark with Jupyter Notebooks

- ▶ We are going to use Jupyter Notebooks for running our walkthroughs & lab exercises.
- ▶ First we need to do the following steps:
 - ▶ Copying all the necessary material in our accounts in Cirrus
 - ▶ Starting an interactive session in a node
 - ▶ Starting a spark cluster (standalone) in that node
 - ▶ Starting a Jupyter session connected with pyspark
- ▶ All the information can be found in “Get_Started_Notebooks_Cirrus”:
https://github.com/EPCCed/prace-spark-for-data-scientists/blob/master/Get_Started_Notebooks_Cirrus.pdf



Submit job via spark-submit

spark-submit Syntax

```
spark-submit --option value \  
  application jar | python file [application arguments]
```

Check the guide - Submitting Spark Applications:

https://github.com/EPCCed/prace-spark-for-data-scientists/blob/master/Spark_Applications/Submitting_Spark_Applications.pdf



Submit job via spark-submit

```
$SPARK_HOME/bin/spark-submit \  
--class <main-class> \  
--master <master-url> \  
--deploy-mode <deploy-mode> \  
--conf \  
....  
<application-jar> [arguments] |  
<python file >[arguments]
```



Some spark-submit options

- ▶ master – Determines how to run the job:
 - ▶ spark://r1i2n5:7077
 - ▶ local
- ▶ driver-memory
 - ▶ amount memory available for the driver process.
- ▶ executor-memory
 - ▶ amount of memory allocated to the executor process
- ▶ executor-cores
 - ▶ total number of cores allocated to the executor process
- ▶ total-executor-cores
 - ▶ Total number of cores available for all executors.

See: <https://spark.apache.org/docs/latest/submitting-applications.html>



Cirrus

- ▶ High-performance computing cluster
- ▶ One of the EPSRC Tier-2 National HPC Services.
- ▶ 280 nodes: 36 Intel Xeon CPUs, hyper threading, 256GB
 - ▶ Each node has (virtually) 72 cores
- ▶ 406 TB of storage – Lustre file system
- ▶ Link: <http://www.cirrus.ac.uk/>

<https://cirrus.readthedocs.io/en/latest/user-guide/connecting.html>



Cirrus

- ▶ Connecting to Cirrus
`ssh [userID]@login.cirrus.ac.uk`
- ▶ Two types of nodes:
 - ▶ Login – access to outside network
 - ▶ Computing – only network between nodes (no access to outside world)
- ▶ For cloning the repository -> use the login node
- ▶ <https://cirrus.readthedocs.io/en/latest/user-guide/connecting.html>



Running jobs on Cirrus

- ▶ PBSPro to schedule jobs
 - ▶ Submission script to submit a job a queue
 - ▶ Interactive jobs are also available
 - ▶ To submit a request for an interactive job reserving 1 nodes (72 physical cores) for 1 hour you would issue the following qsub command from the command line

```
qsub -IVl select=3:ncpus=36,walltime=01:00:00,place=scatter:excl -A y15 -  
q <reservation number> -j oe
```

- ▶ Your session will end:
 - ▶ It hits the requested walltime
 - ▶ Typing exit command within the session
- ▶ <https://cirrus.readthedocs.io/en/latest/user-guide/batch.html#interactive-jobs>



Jupyter notebooks

- ▶ Start the jupyter server:

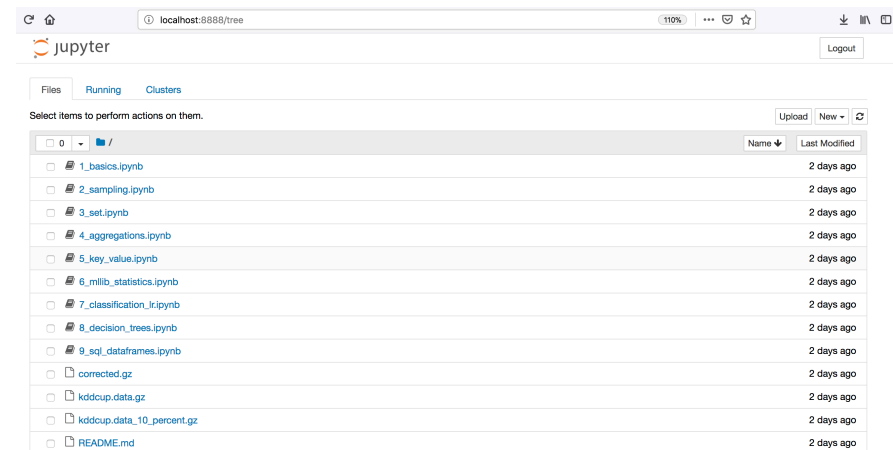
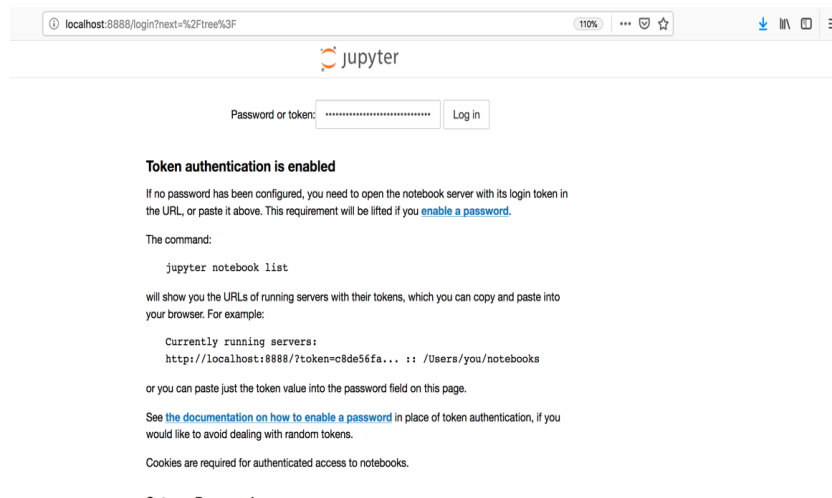
- ▶ `./start_Jupyter_local.sh` will give you a token, like this one:

`http://0.0.0.0:8888/?token=2d5e554b2397355c334b8c3367503b06c4f6f95a26151795`

- ▶ Open another terminal and type the following command

`ssh USER@login.cirrus.ac.uk -L8888:MASTER NODE:8888`


- ▶ Got to a Web browser at `http://localhost:8888`



All the information can be found at “Get_Started_Notebooks_Cirrus”:
https://github.com/EPCCed/prace-spark-for-data-scientists/blob/master/Get_Started_Notebooks_Cirrus.pdf



Master Spark UI

 **Spark Master at spark://r1i1n20:7077**

URL: spark://r1i1n20:7077
Alive Workers: 1
Cores in use: 72 Total, 0 Used
Memory in use: 250.6 GB Total, 0.0 B Used
Applications: 0 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

▼ Workers (1)


Worker Id	Address	State	Cores	Memory
worker-20190106070903-10.148.0.44-32960	10.148.0.44:32960	ALIVE	72 (0 Used)	250.6 GB (0.0 B Used)

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

▼ Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------



Drive Spark Jobs (?)

User: rosaf3

Total Uptime: 3.5 min

Scheduling Mode: FIFO

Completed Jobs: 10

▶ Event Timeline

▼ Completed Jobs (10)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
9	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:36:33	71 ms	1/1	4/4
8	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:36:33	61 ms	1/1	1/1
7	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:34:31	64 ms	1/1	4/4
6	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:34:31	60 ms	1/1	1/1
5	count at <ipython-input-7-bcf6b48ba43a>:1 count at <ipython-input-7-bcf6b48ba43a>:1	2019/01/06 10:34:28	0.3 s	1/1	72/72
4	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:34:23	24 ms	1/1	1/1
3	count at <ipython-input-4-e13515b0683a>:4 count at <ipython-input-4-e13515b0683a>:4	2019/01/06 10:34:19	0.9 s	1/1	1/1
2	runJob at PythonRDD.scala:153 runJob at PythonRDD.scala:153	2019/01/06 10:34:19	31 ms	1/1	1/1

Every SparkContext launches a web UI (Spark driver's web UI), by default on port 4040, that displays useful information about the application.

ssh USER@login.cirrus.ac.uk -L4040:DRIVER NODE:4040


web browser → localhost:4040



Running notebooks in your laptop

- ▶ **Prerequisites: Anaconda, Python3**
- ▶ Get Spark from the [downloads page](#) of the project website
(<https://blog.sicara.com/get-started-pyspark-jupyter-guide-tutorial-ae47e847c94f>)
- ▶ Check if pyspark is properly install → type pyspark in a terminal

Welcome to



```
>> git clone https://github.com/EPCCed/prace-spark-for-data-scientists.git
>> cd walkthrough_examples
>> export SPARK_HOME=[INSTALLATION_PATH]/spark-2.4.0-bin-hadoop2.7/
>> PYSPARK_DRIVER_PYTHON=jupyter PYSPARK_DRIVER_PYTHON_OPTS='notebook' \
$SPARK_HOME/bin/pyspark
```




PARTNERSHIP FOR ADVANCED COMPUTING IN EUROPE

THANK YOU FOR YOUR ATTENTION

www.prace-ri.eu