


SPARK CLUSTER OVERVIEW

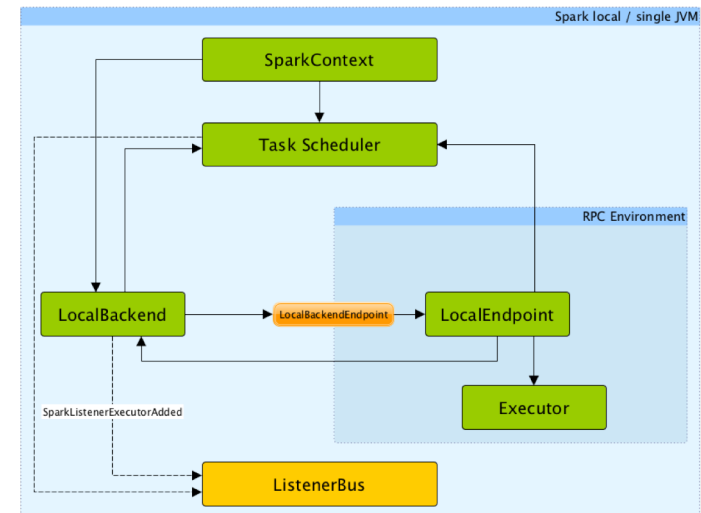
Spark Execution modes

It is possible to run a spark application using **cluster mode**, **local mode** (pseudo-cluster) or with an **interactive** shell (*pyspark* or *spark-shell*).

- Cluster mode

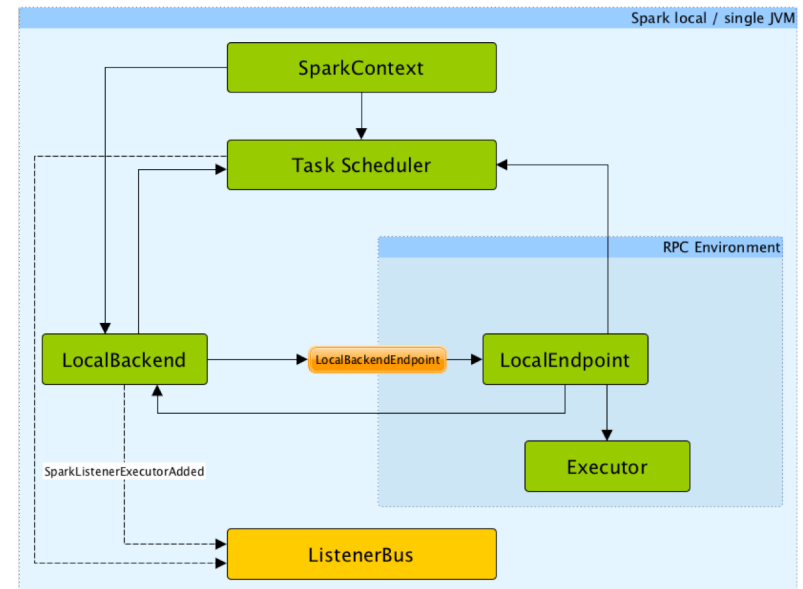


- Local mode



Spark Execution – Local mode


- In this non-distributed single-JVM deployment mode.
- Spark spawns all the execution components - [driver](#), [executor](#), [LocalSchedulerBackend](#), and [master](#) - in the same single JVM.
- The default parallelism is the number of threads as specified in the [master URL](#).
- This is the only mode where a driver is used for execution



Spark Execution – Cluster mode

It currently provides several options:


- **Standalone Deploy Mode**: simplest way to deploy Spark on a private cluster
- Apache Mesos
- Hadoop YARN
- Kubernetes



Spark is agnostic to the underlying cluster manager

Spark Execution – Cluster mode

- Spark applications are run as independent sets of processes, coordinated by a SparkContext in a (*) *driver* program.
- The *context* connects to the cluster manager *which allocates resources*.




- Each *worker* in the cluster is managed by an *executor*.
- The *executor* manages computation as well as storage and caching on each machine.

(*) *driver* → process running the *main()* function of the application and creating the *SparkContext*

Spark Execution – Cluster mode

- The application code is sent from the *driver* to the *executors*, and the executors specify the context and the various *tasks* to be run.
- The *driver* program must listen for and accept incoming connections from its executors throughout its lifetime.





Spark Execution – Cluster mode

Spark App

Each SparkContext creates a Spark application, which includes a lot of scheduling components.

Upon an **Action**, the driver program submits the job to the cluster manager.



Cluster manager





Start executors on Worker Nodes.

It does **not** know about stages.

Worker

Launch Spark Executor in a process.

Tasks are launched in separate threads, one per each core on the worker node (can be configured)



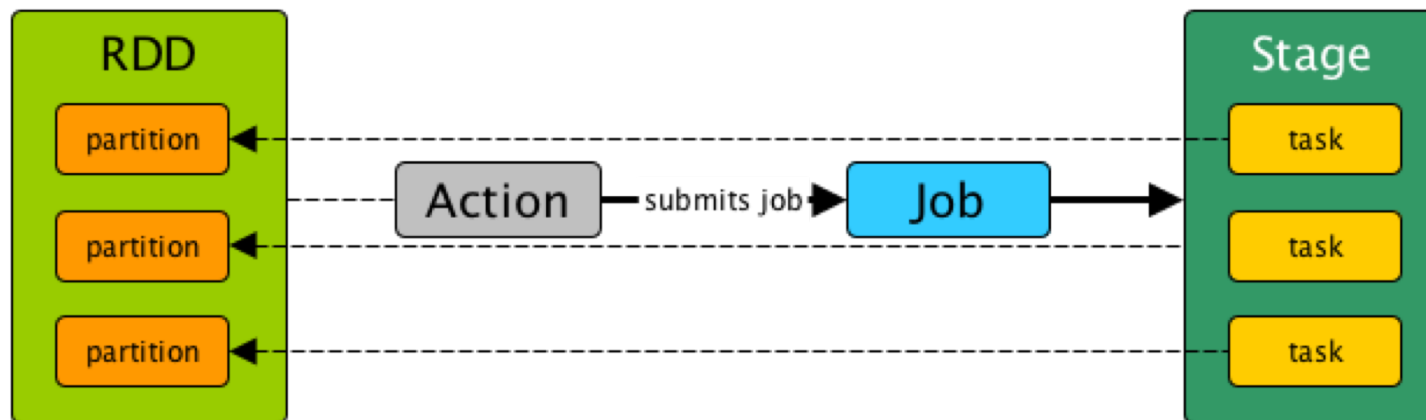
Spark – Standalone Cluster – Deploy modes

For standalone clusters supports two deploy modes.
They distinguish where the *driver* process runs:

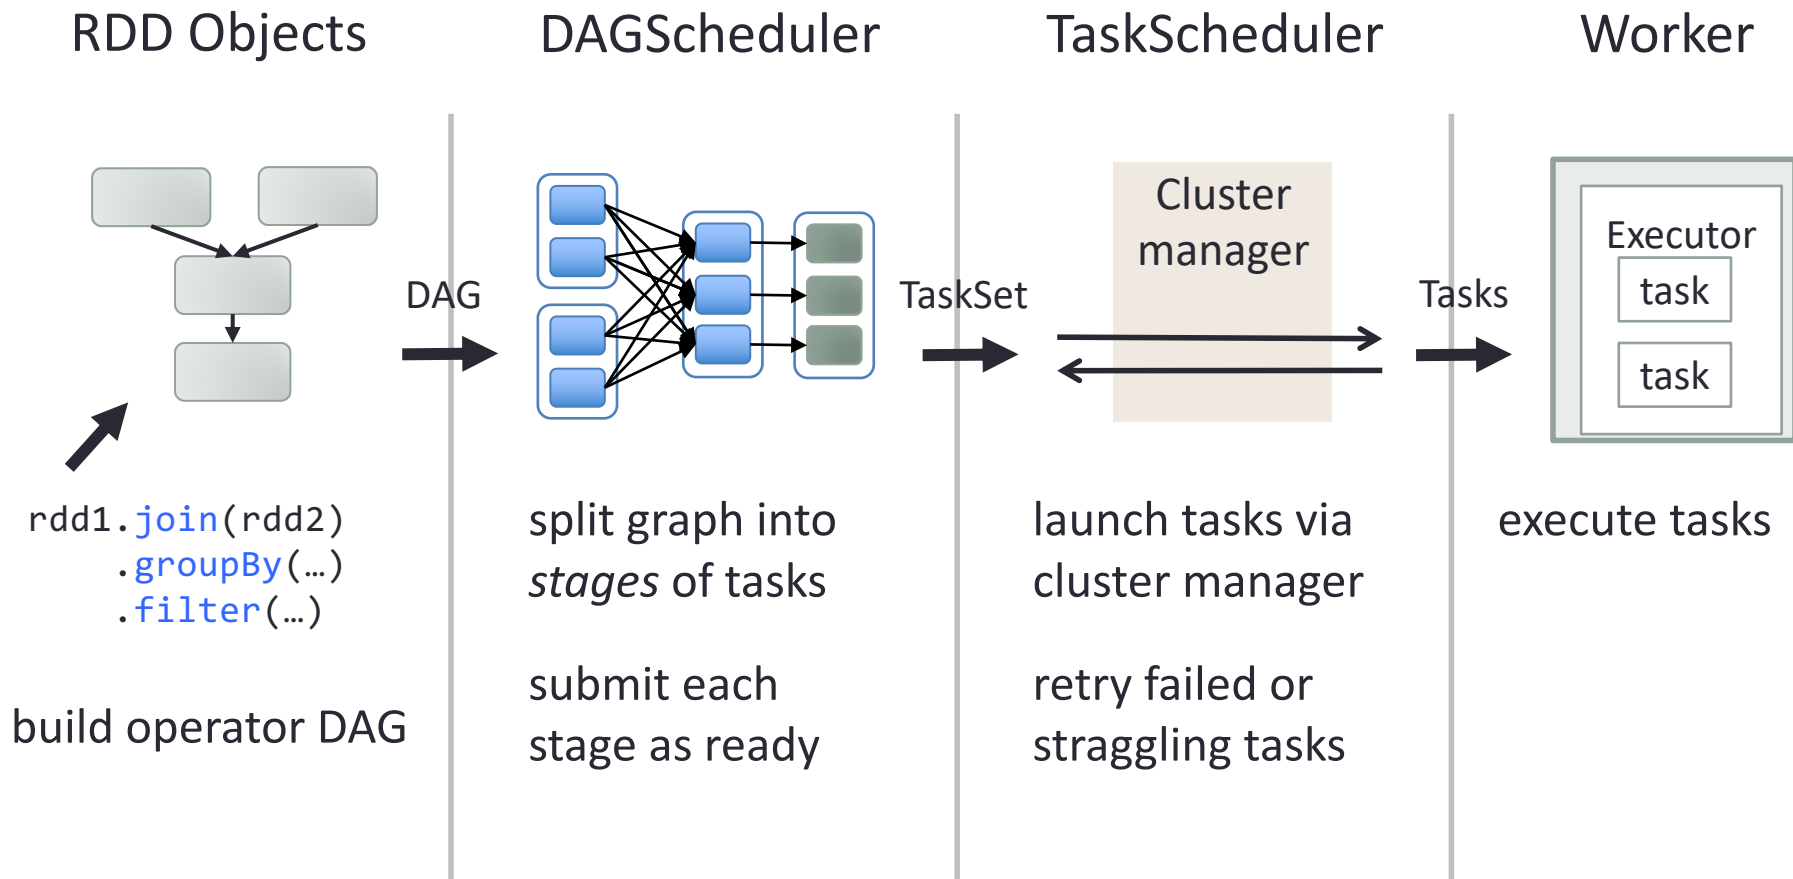
- *Client mode*: the *driver* is launched in the *same process as the client* that submits the application.
- *Cluster mode*: the *driver* is launched from *one of the Worker processes* inside the cluster.
 - The client process exits as soon as it fulfils its responsibility of submitting the application without waiting for the application to finish.

Spark Components

- Task: individual unit of work sent to one executor over a sequences of partitions
- Job : set of tasks executed as a result of an action
- Stage: set of tasks in a job that can be executed in parallel
- RDD: Parallel dataset with partitions
- DAG: Logical Graph of RDD operations



Job scheduling



Spark Application – wordcount.py


The application that we are going to create is a simple “wordcount”:

- Performs a ***textFile*** operation to read an input file in HDFS
- ***flatMap*** operation to split each line into words
- ***map*** operation to form (word, 1) pairs
- ***reduceByKey*** operation to sum the counts for each word.

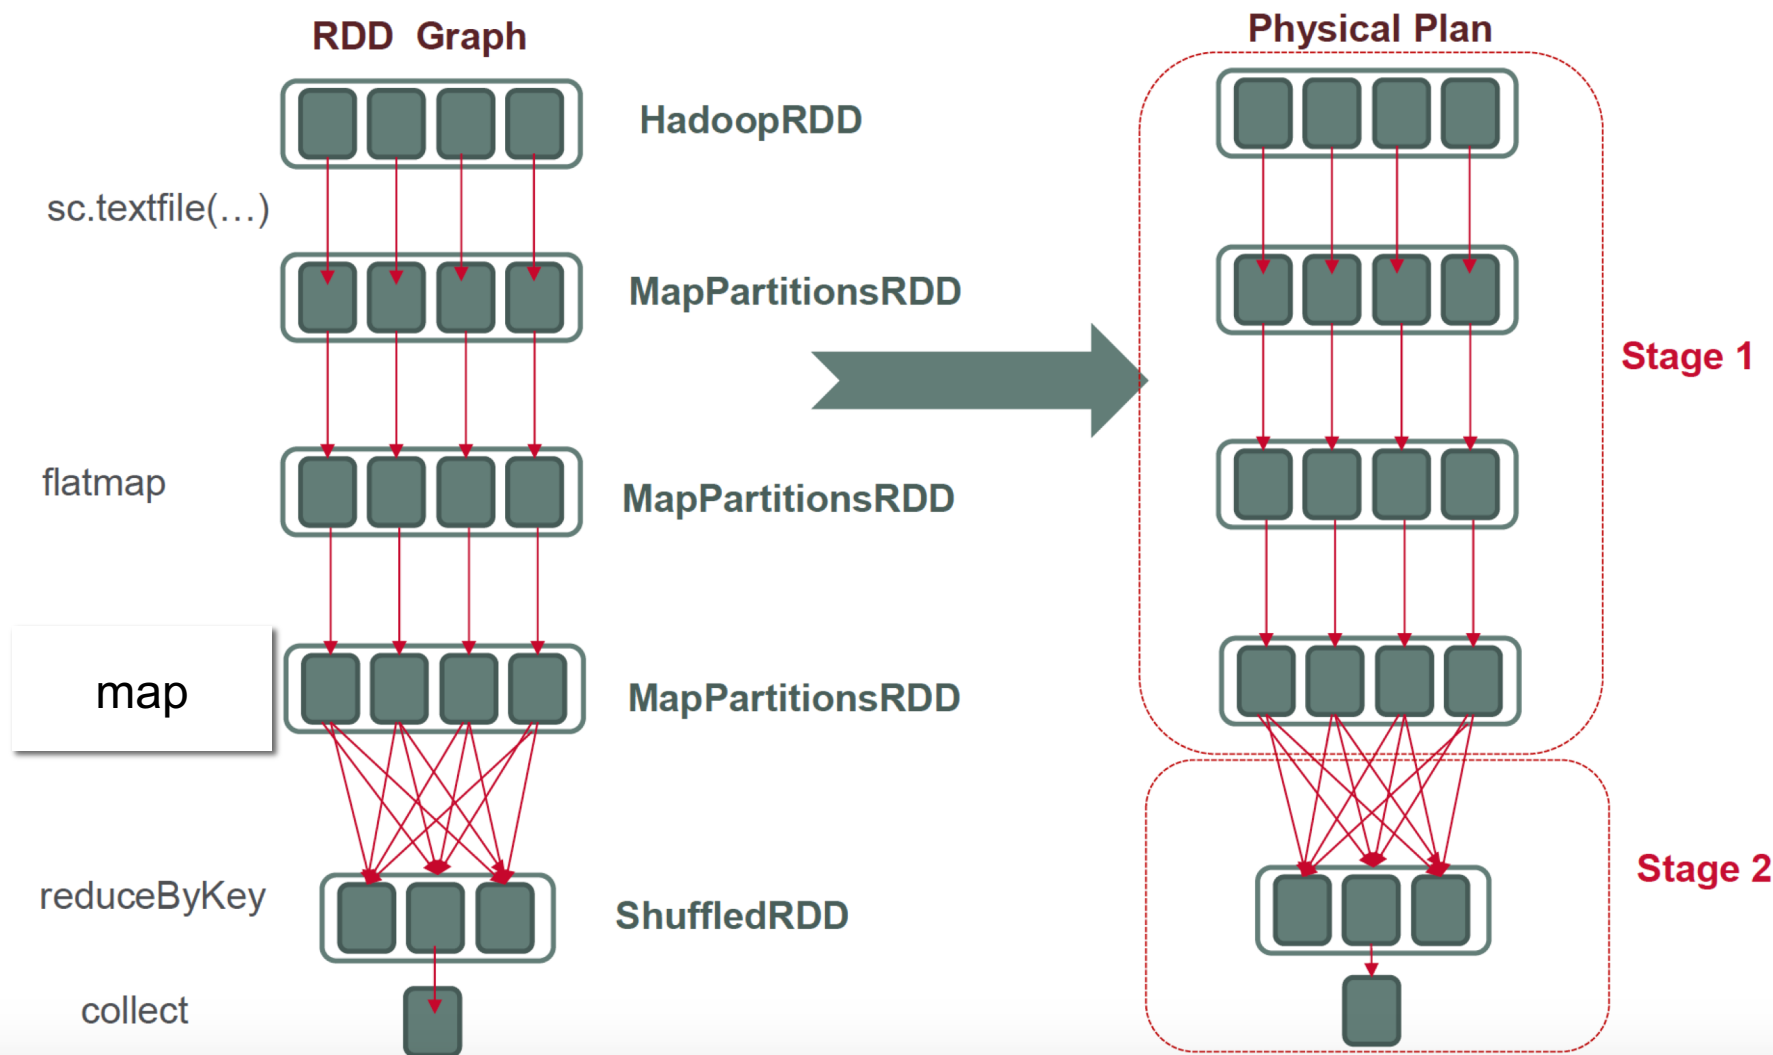
Spark Application – wordcount.py

```
import sys
from pyspark import SparkContext, SparkConf

if __name__ == "__main__":
    conf = SparkConf().setAppName("Spark Count")
    sc = SparkContext(conf=conf)
    logFile = "../spark-2.4.0-bin-hadoop2.7/README.md"
    textFile = sc.textFile(logFile)
    wordCounts = textFile.flatMap(lambda line: line.split()).\
        map(lambda word: (word, 1)).reduceByKey(lambda a, b: a+b)
    output=wordCounts.collect()
    for (word, count) in output:
        print("%s: %i" % (word, count))
```

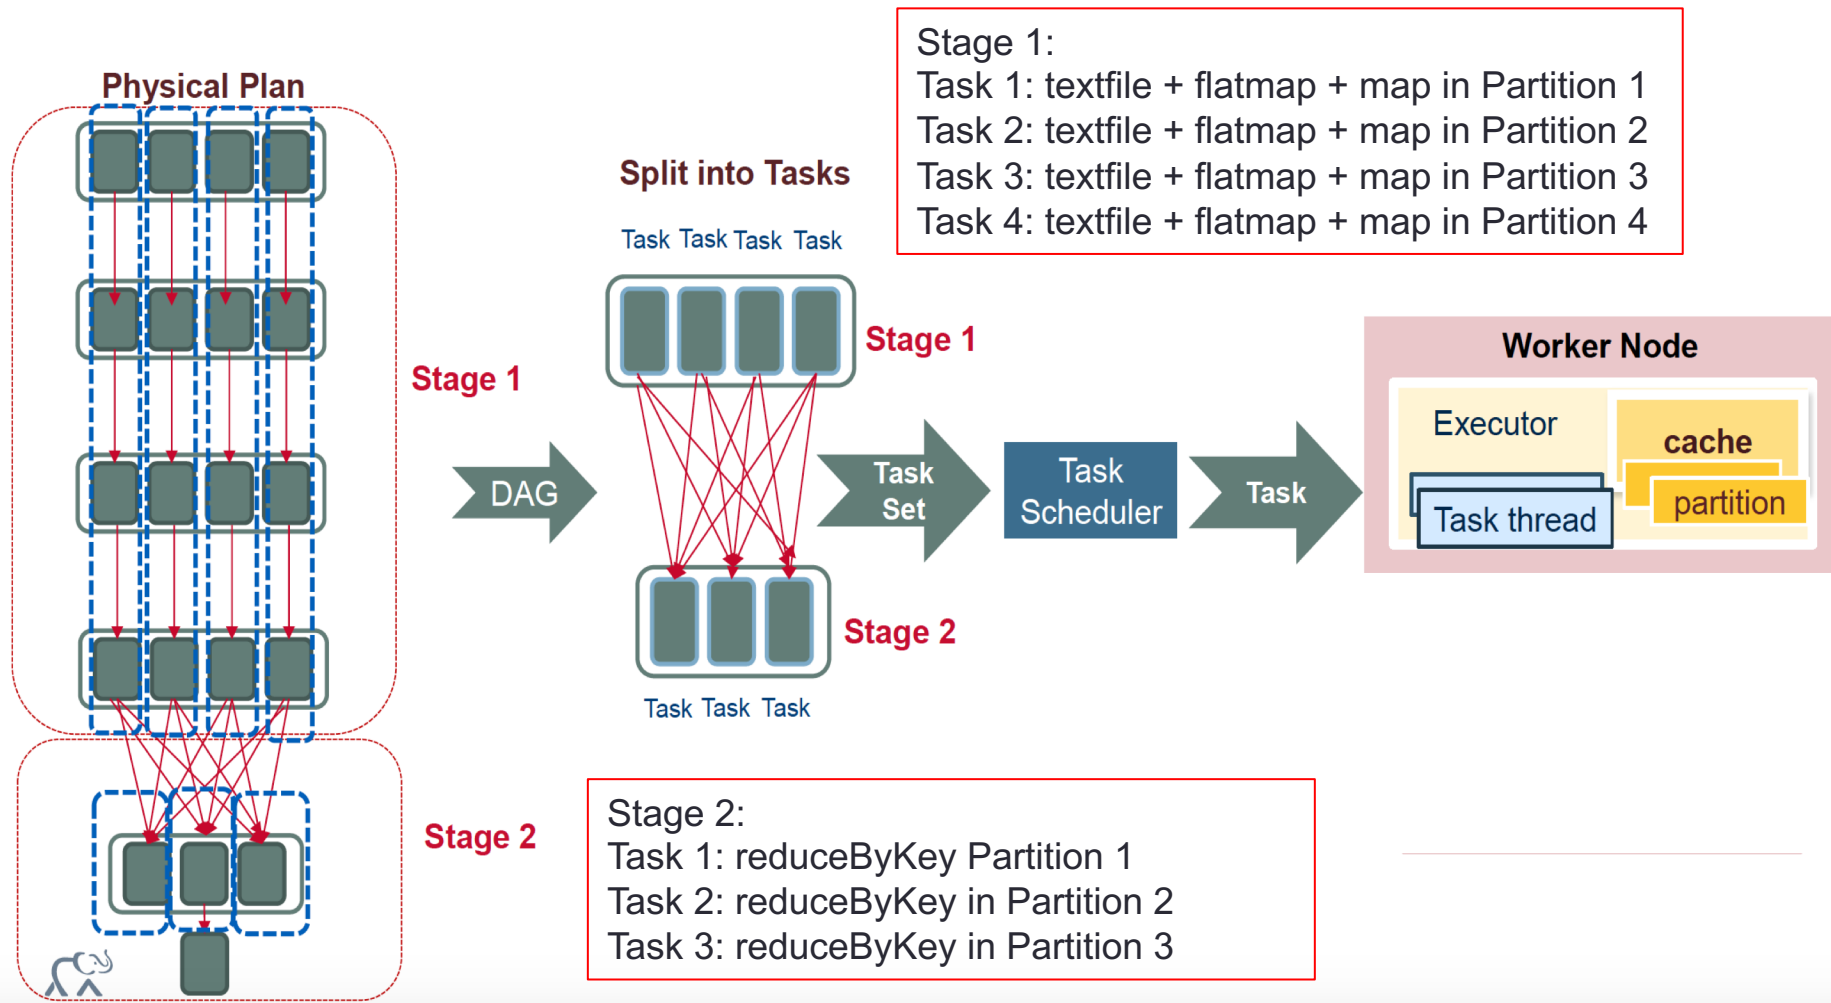


RDD DAG -> Physical Execution plan



Initial RDD distributed among 4 partitions. Final RDD distributed among 3 partitions

Execution plan -> Stages and Tasks



Operations that can run on the same partition are executed in stages

Running Spark Applications

- **Notebooks** are great for:
 - developing and testing quickly experiment with the data
 - demos and collaborating with other people
- **Spark-submit** jobs are more likely to be used in **production**.

Running Spark with Jupyter Notebooks

We are going to use Jupyter Notebooks for running our walkthroughs & lab exercises.

First we need to do the following steps:

- Copying all the material necessaire in our accounts in Cirrus
- Starting an interactive session in a node
- Starting a spark cluster (standalone) in that node
- Starting a Jupyter session connected with pyspark

All the information can be found at “Get_Started_Notebooks_Cirrus”:

https://github.com/EPCCed/prace-spark-for-data-scientists/blob/master/Get_Started_Notebooks_Cirrus.pdf

Submit job via spark-submit

spark-submit Syntax

```
spark-submit --option value \  
  application jar | python file [application arguments]
```

Check the guide - Submitting Spark Applications:

https://github.com/EPCCed/prace-spark-for-data-scientists/blob/master/Spark_Applications/Submitting_Spark_Applications.pdf

Submit job via spark-submit

```
$SPARK_HOME/bin/spark-submit \  
--class <main-class> \  
--master <master-url> \  
--deploy-mode <deploy-mode> \  
--conf \  
....  
<application-jar> [arguments] |  
<python file >[arguments]
```

Some spark-submit options

- master – Determines how to run the job:
 - spark://r1i2n5:7077
 - local
- driver-memory
 - amount memory available for the driver process.
- executor-memory
 - amount of memory allocated to the executor process
- executor-cores
 - total number of cores allocated to the executor process
- total-executor-cores
 - Total number of cores available for all executors.

Cirrus

- High-performance computing cluster
- One of the EPSRC Tier-2 National HPC Services.
- 280 nodes: 36 Intel Xeon CPUs, hyper threading, 256GB
- 406 TB of storage- Lustre
- Link: <http://www.cirrus.ac.uk/>
-